

# *Feature Detection & Invariance*

<Vision System>

Department of Robot Engineering  
Prof. Younggun Cho

Most of these slides were adapted from:

- Kris Kitani (15-463, Fall 2016).

Some slides were inspired or taken from:

- Fredo Durand (MIT).
- James Hays (Georgia Tech).
- Ionannis Gkioulekas (CMU).

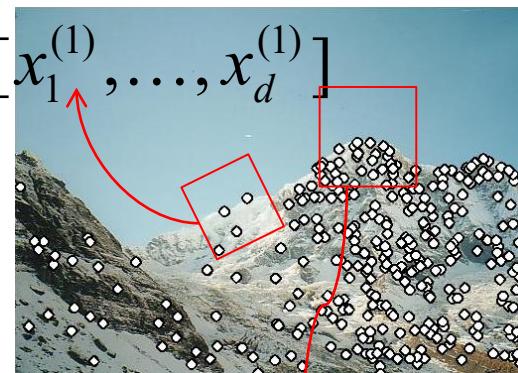


# Local features: main components

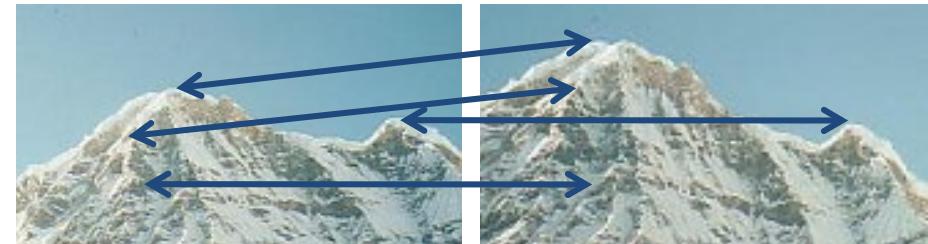
- 1) Detection: Identify the interest points



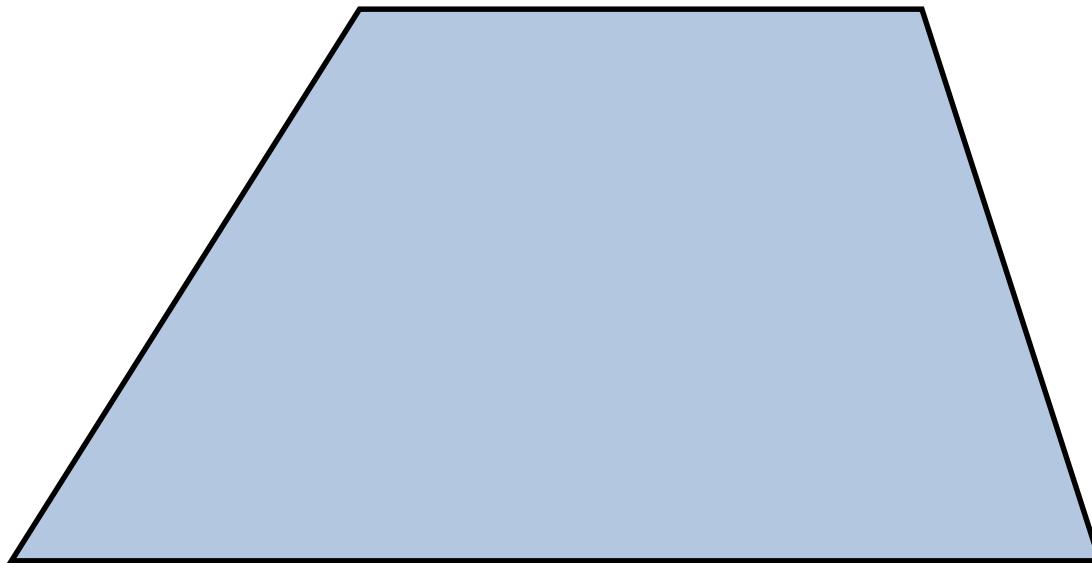
- 2) Description: Extract vector feature descriptor surrounding  $\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$  each interest point.



- 3) Matching: Determine correspondence between descriptors in two views



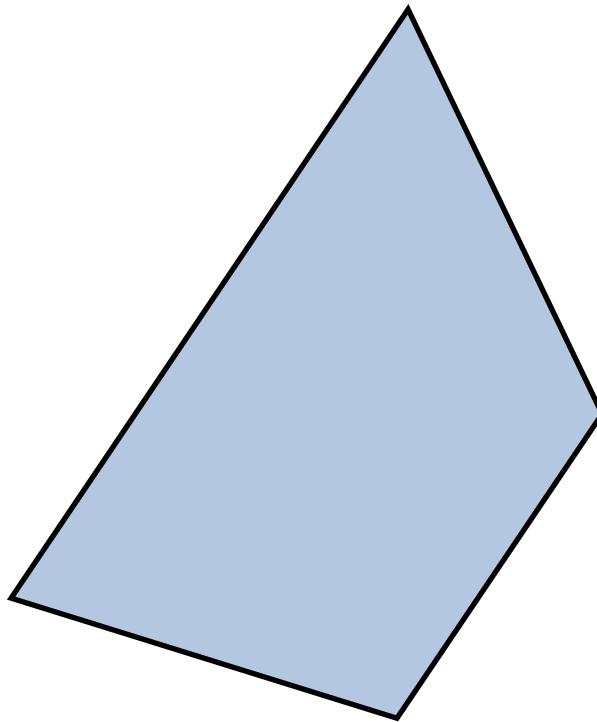
# Why feature?



Pick a point in the image.  
Find it again in the next image.

*What type of feature would you select?*

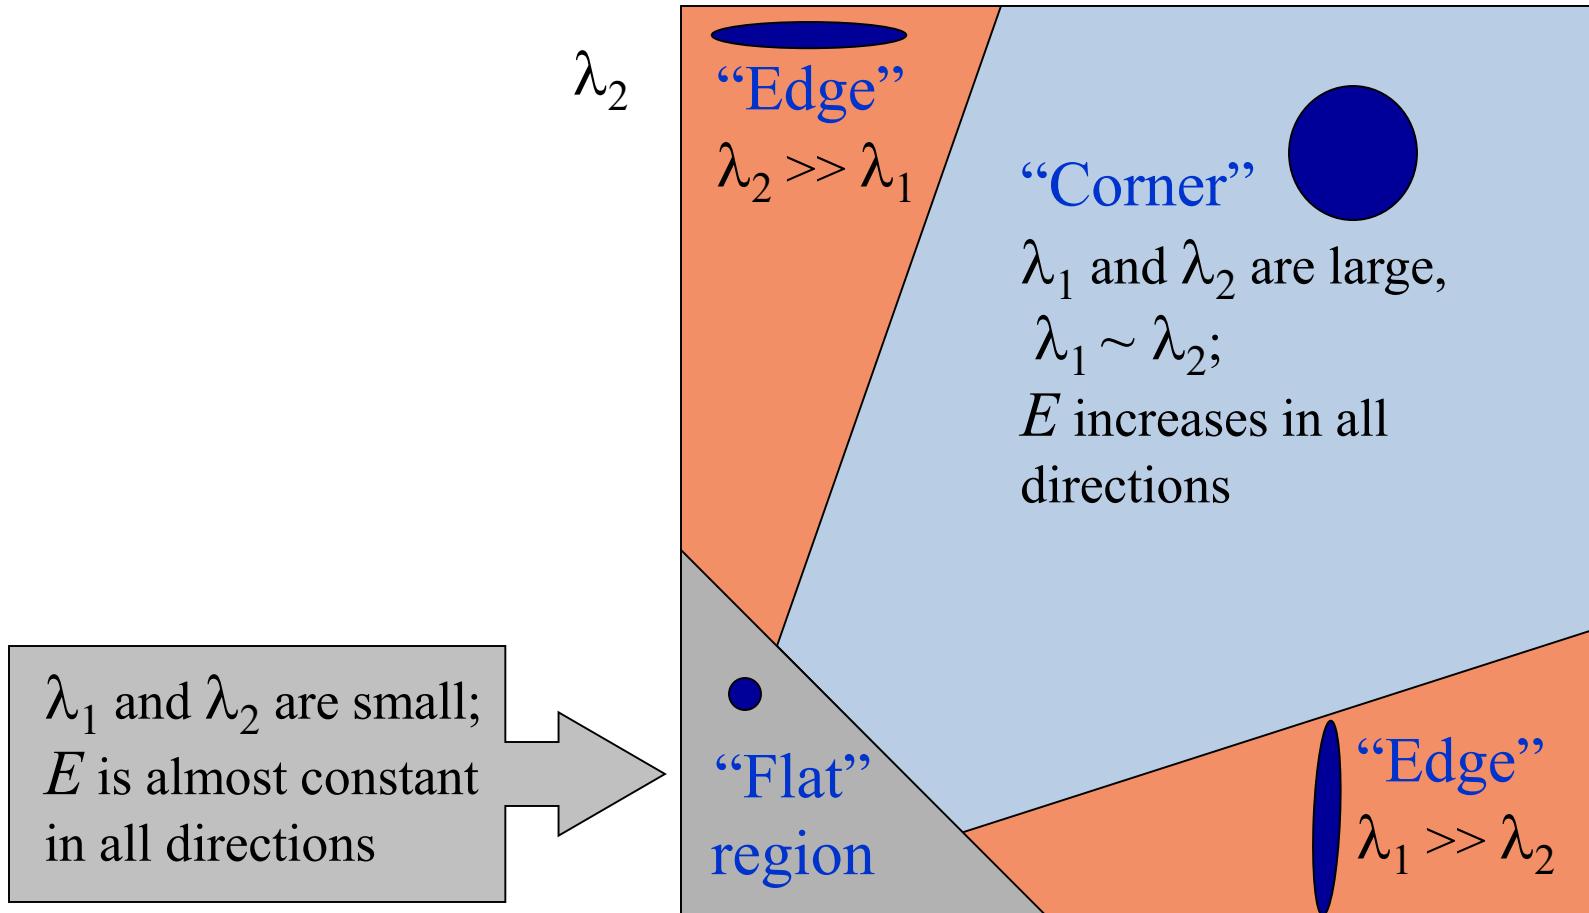
# Why feature?



Pick a point in the image.  
Find it again in the next image.

*What type of feature would you select?*

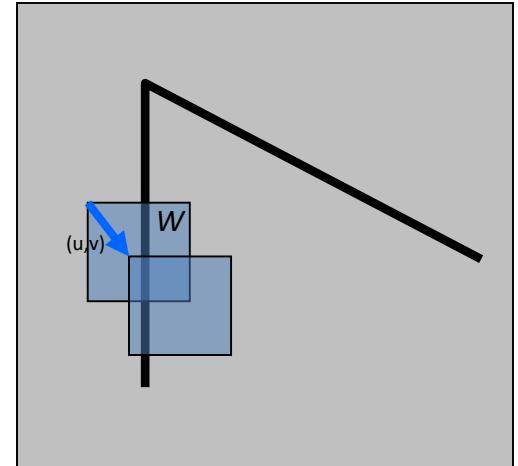
# What we studied last week.



# Corner detection: the math

Consider shifting the window  $W$  by  $(u, v)$

- define an SSD “error”  $E(u, v)$ :



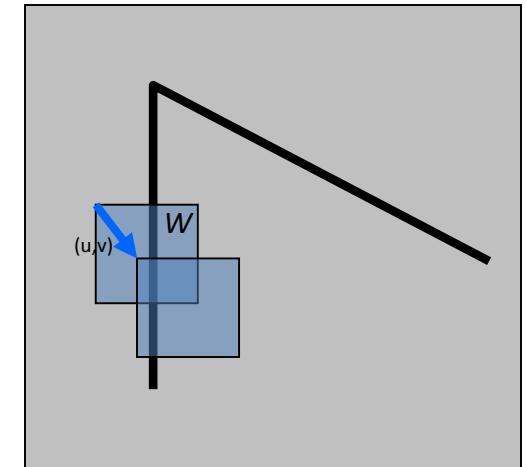
$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \end{aligned}$$

# Corner detection: the math

Consider shifting the window  $W$  by  $(u, v)$

- define an SSD “error”  $E(u, v)$ :

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \\ &\approx A u^2 + 2Buv + Cv^2 \end{aligned}$$



$$A = \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2$$

- Thus,  $E(u, v)$  is locally approximated as a quadratic error function

We will need this to understand the...

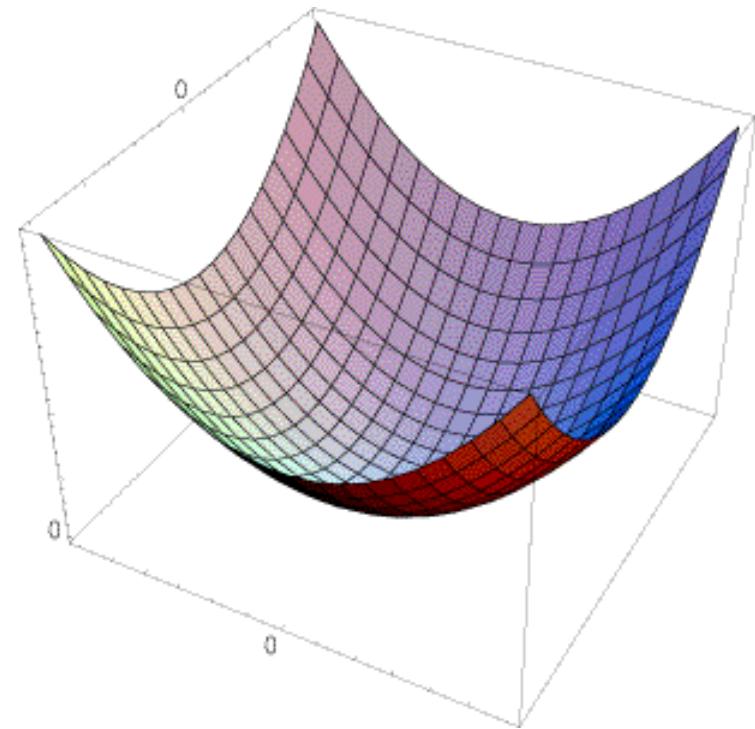
# Error function for Harris Corners

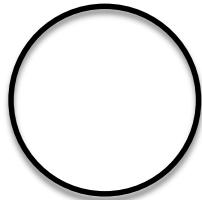
The surface  $E(u,v)$  is locally approximated by a quadratic form

$$E(u,v) \approx [u \ v] H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$H = \sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

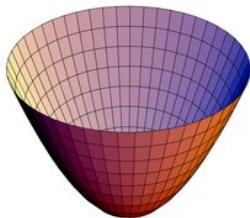
Quadratic Form





Equation of a circle

$$1 = x^2 + y^2$$



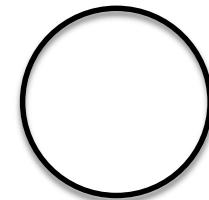
Equation of a ‘bowl’ (paraboloid)

$$f(x, y) = x^2 + y^2$$

*If you slice the bowl at*

$$f(x, y) = 1$$

*what do you get?*



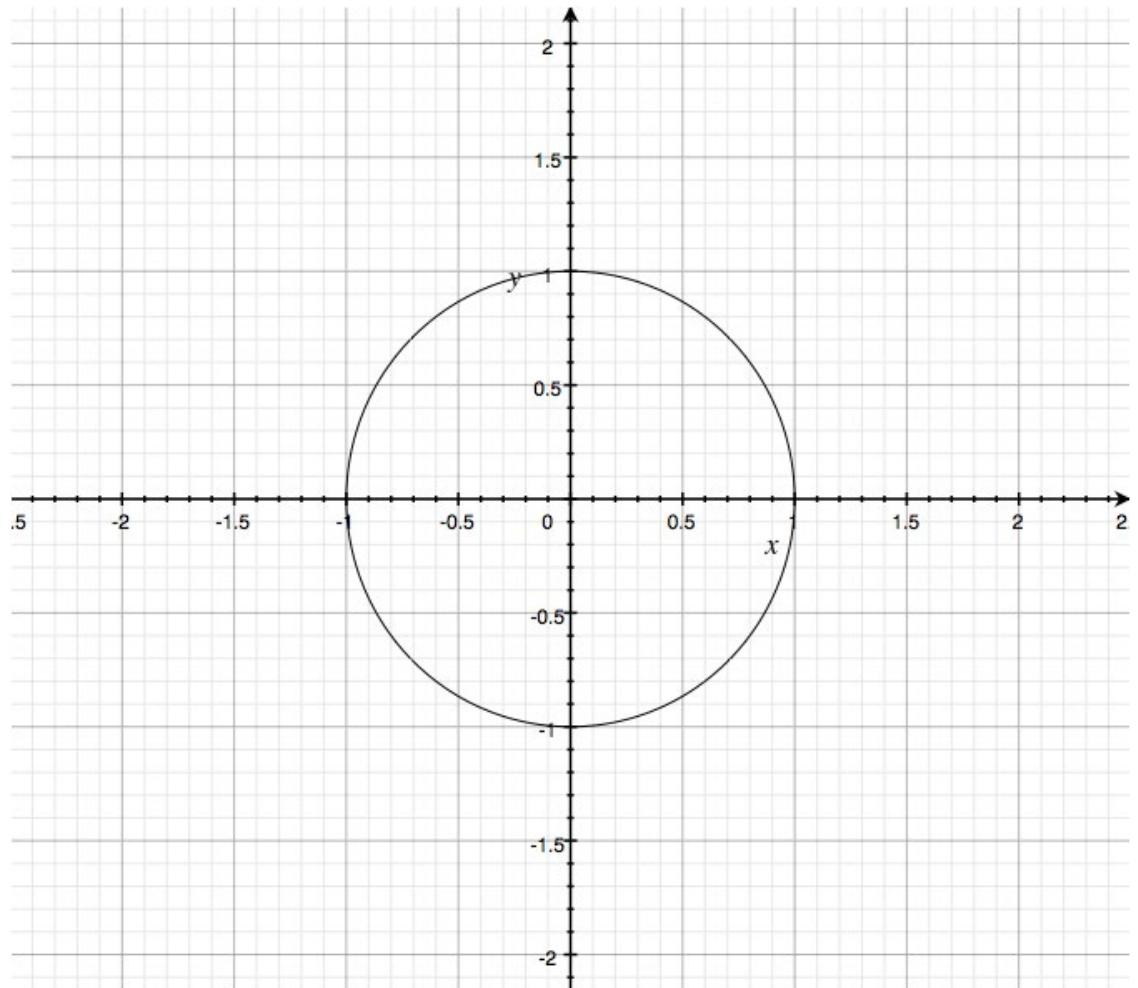
$$f(x, y) = x^2 + y^2$$

can be written in matrix form like this...

$$f(x, y) = \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

$$f(x, y) = \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

'sliced at 1'

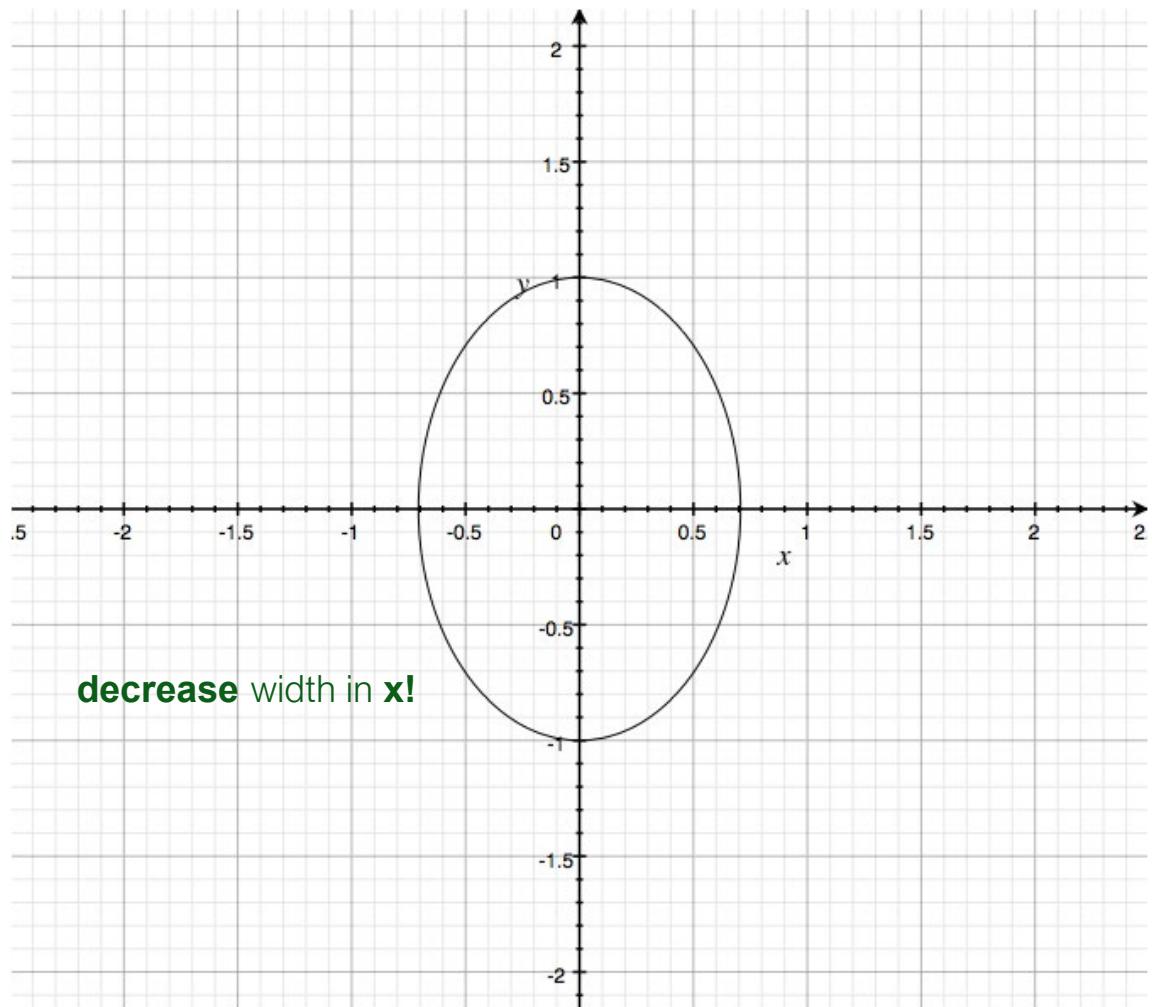


What happens if you **increase** coefficient on **x**?

$$f(x, y) = \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

and slice at 1

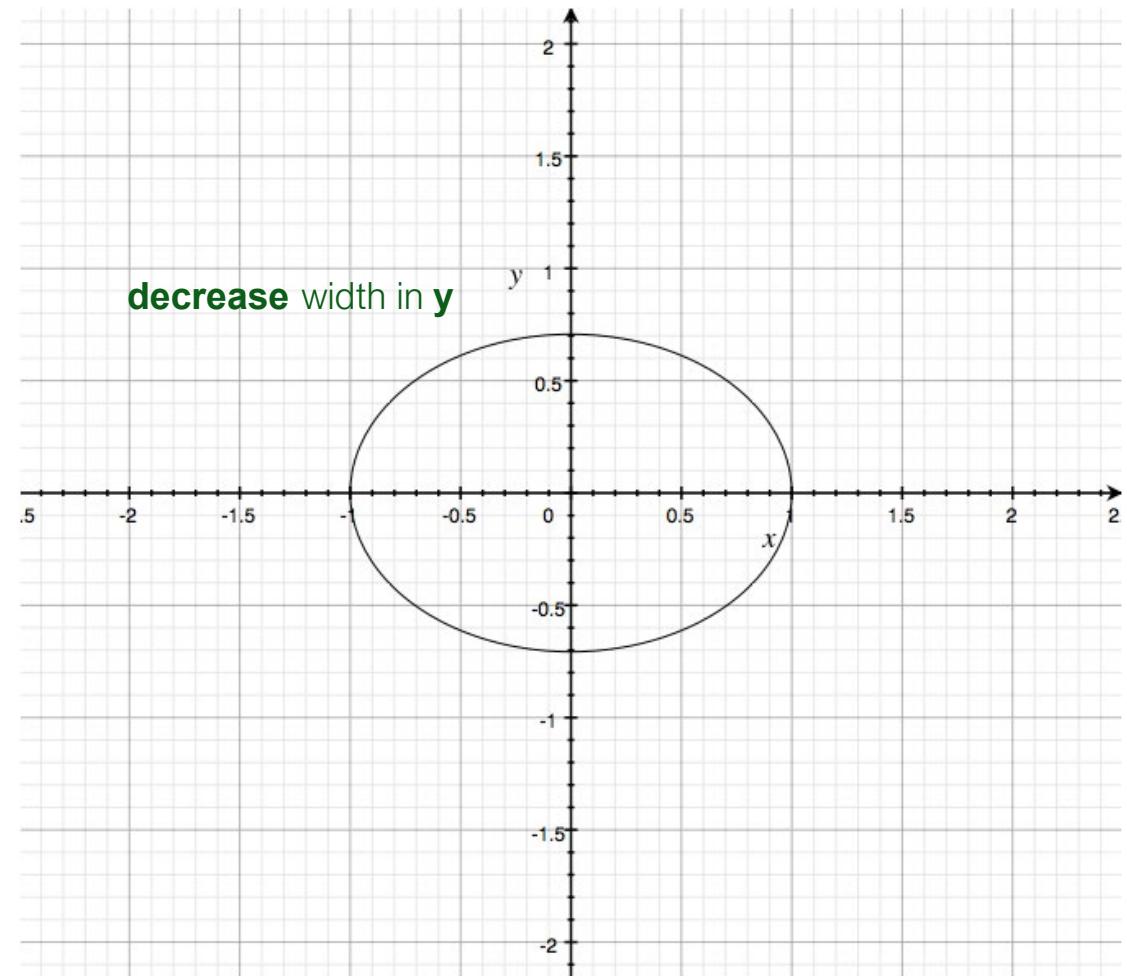
**decrease** width in **x**!



What happens if you **increase** coefficient on **y**?

$$f(x, y) = \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

and slice at 1



$$f(x, y) = x^2 + y^2$$

can be written in matrix form like this...

$$f(x, y) = \begin{bmatrix} x & y \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

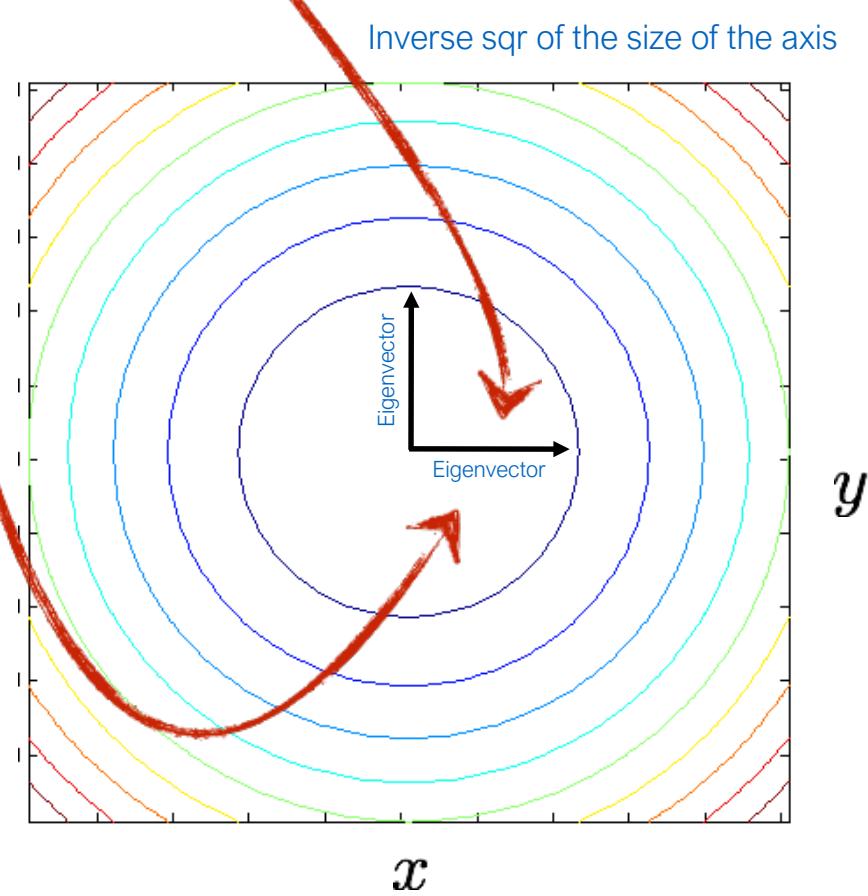
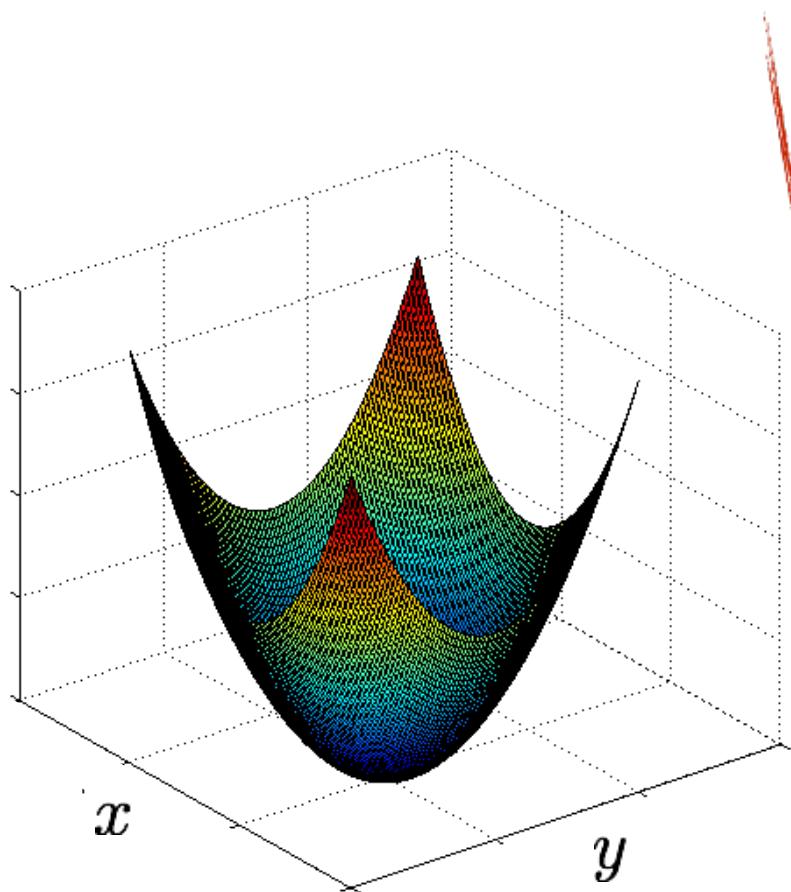
### Result of Singular Value Decomposition (SVD)

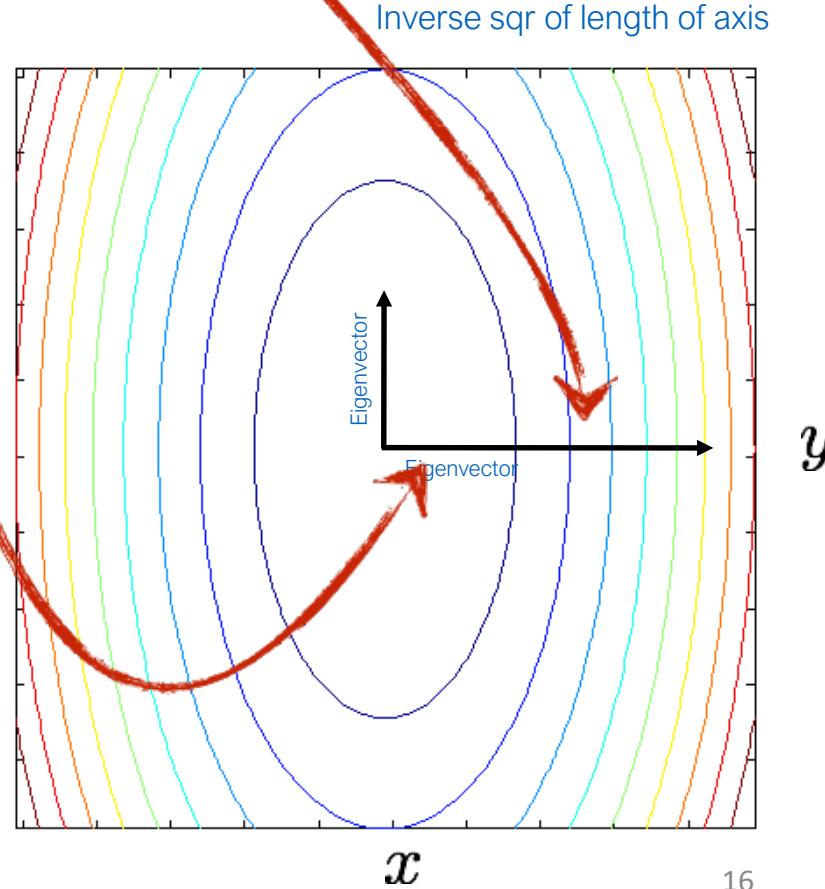
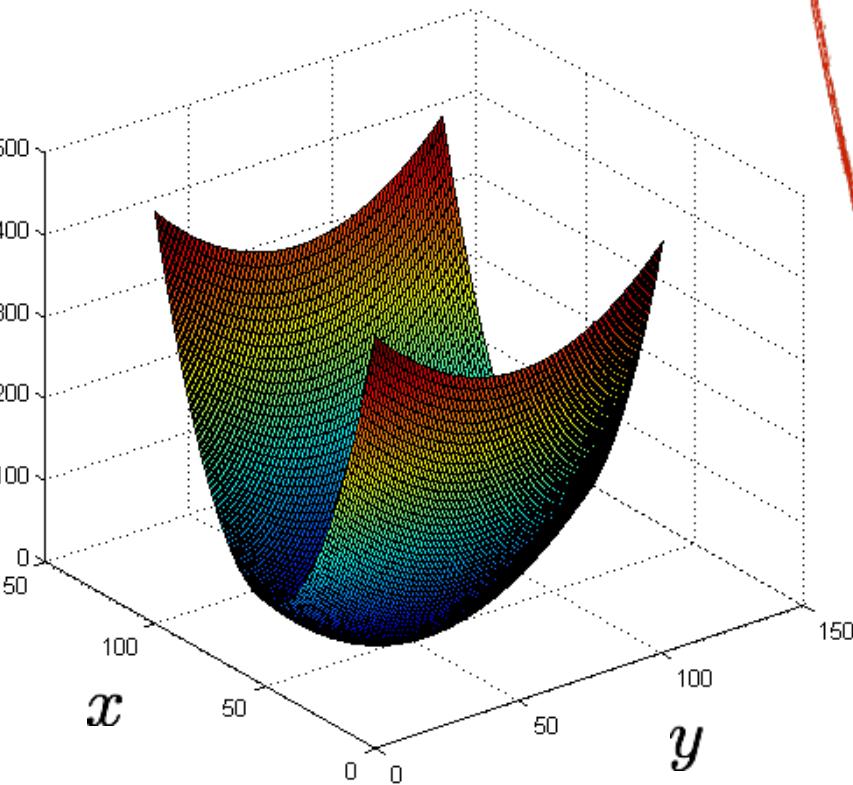
$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} \text{eigenvectors} \\ \text{axis of the 'ellipse slice'} \end{bmatrix} \begin{bmatrix} \text{eigenvalues along diagonal} \\ \text{Inverse sqrt of length of the quadratic along the axis} \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}^\top$$

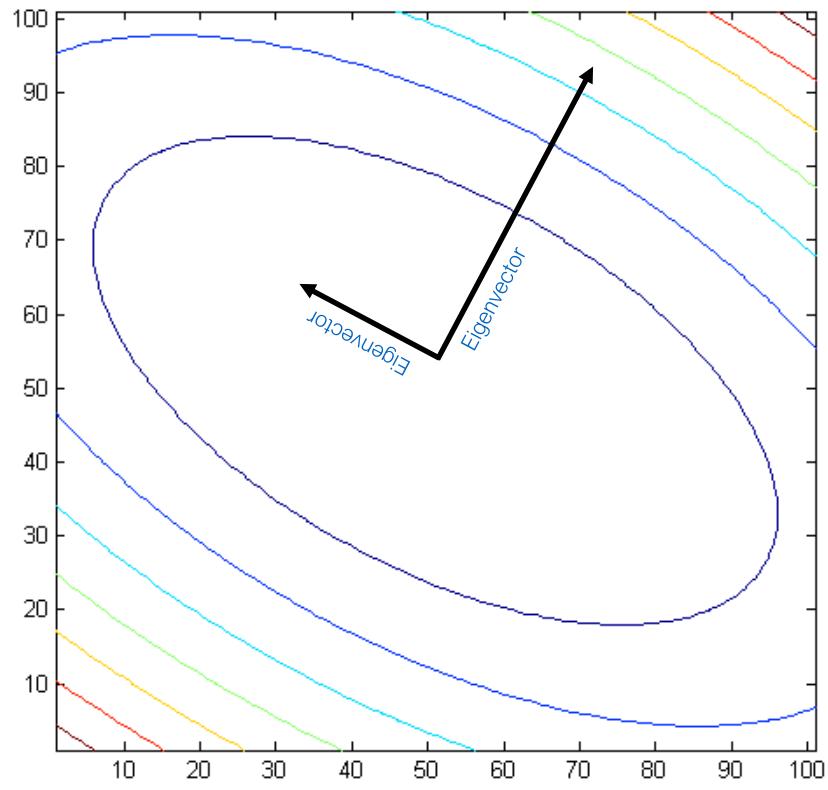
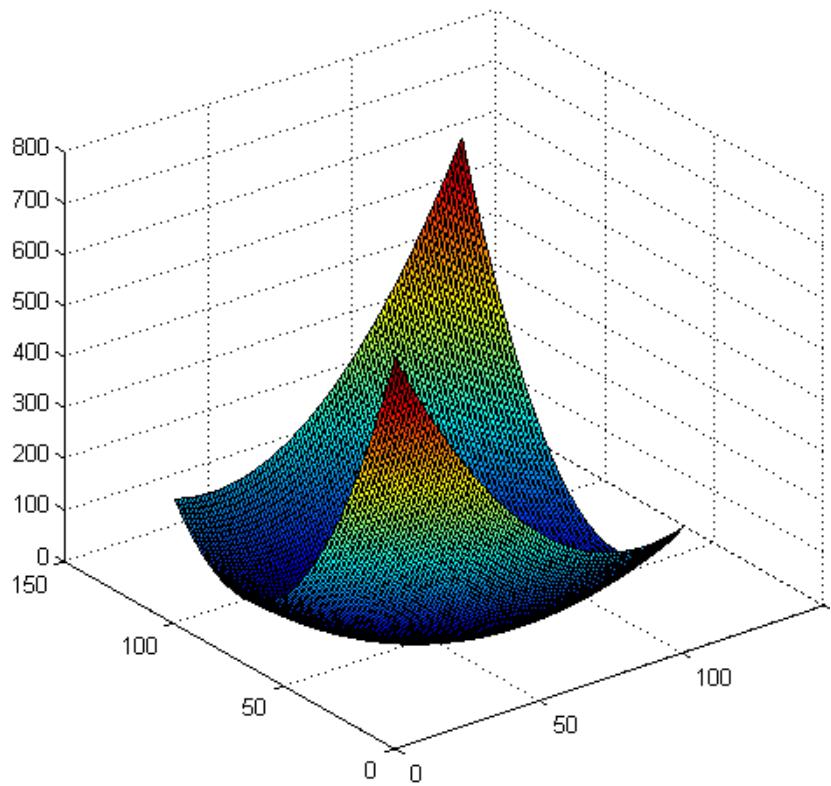
The equation shows the decomposition of the identity matrix into three components. The first component is labeled "eigenvectors" and "axis of the 'ellipse slice'". The second component is labeled "eigenvalues along diagonal" and "Inverse sqrt of length of the quadratic along the axis". The third component is the transpose of the second component.

$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}^T$$

Eigenvectors   Eigenvalues

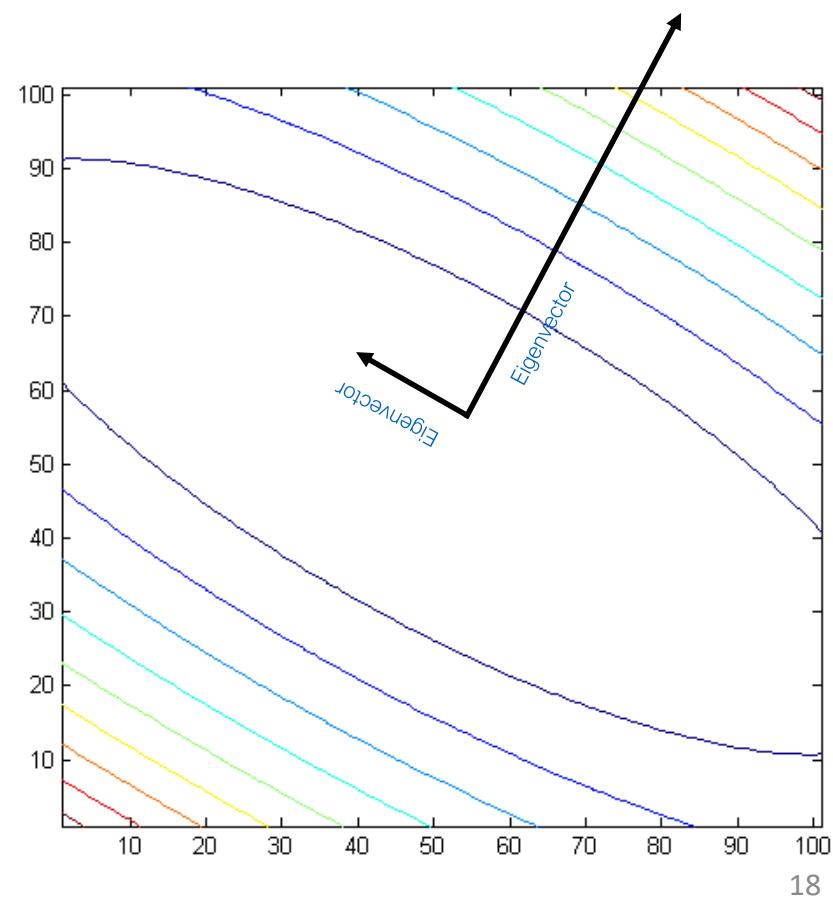
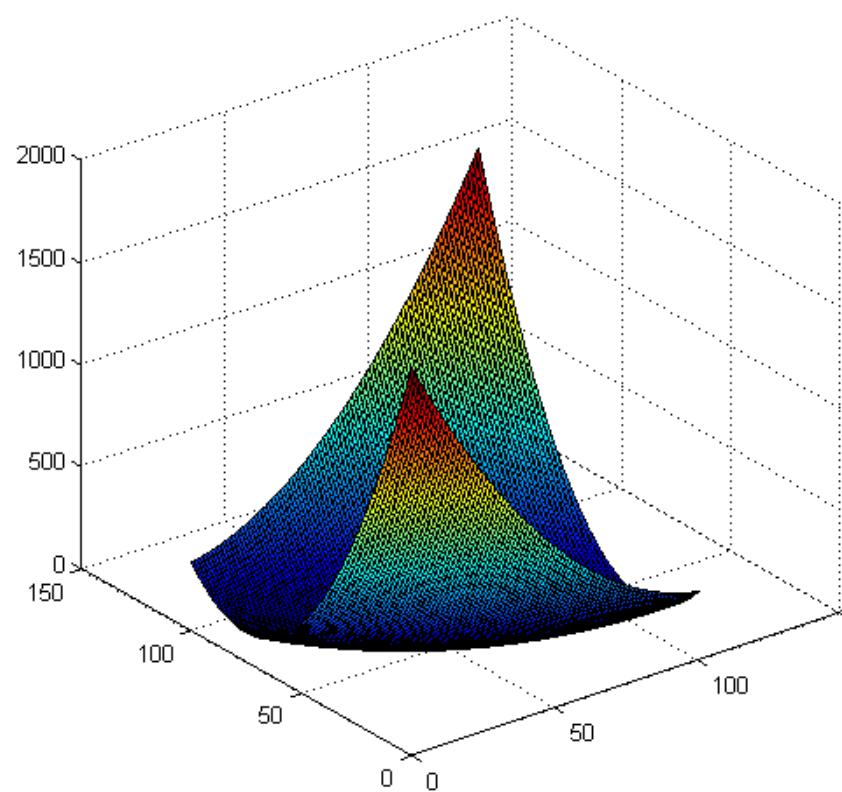




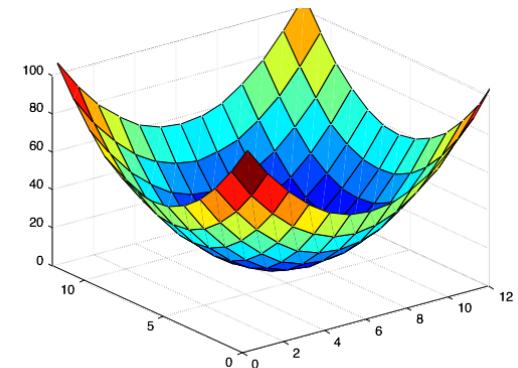
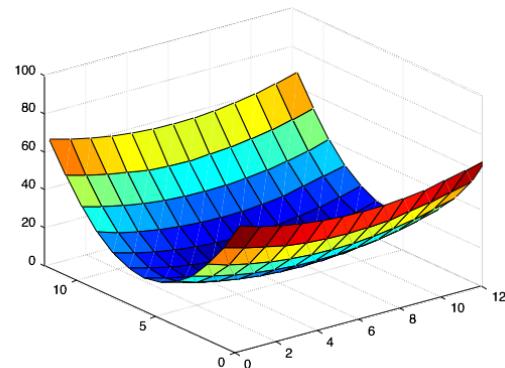
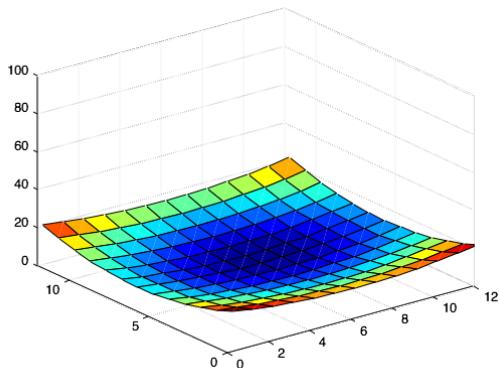


$$A = \begin{bmatrix} 7.75 & 3.90 \\ 3.90 & 3.25 \end{bmatrix} = \begin{bmatrix} 0.50 & -0.87 \\ -0.87 & -0.50 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 10 \end{bmatrix} \begin{bmatrix} 0.50 & -0.87 \\ -0.87 & -0.50 \end{bmatrix}^T$$

Eigenvalues  
Eigenvectors      Eigenvalues  
Eigenvectors



*Which error surface indicates a good image feature?*



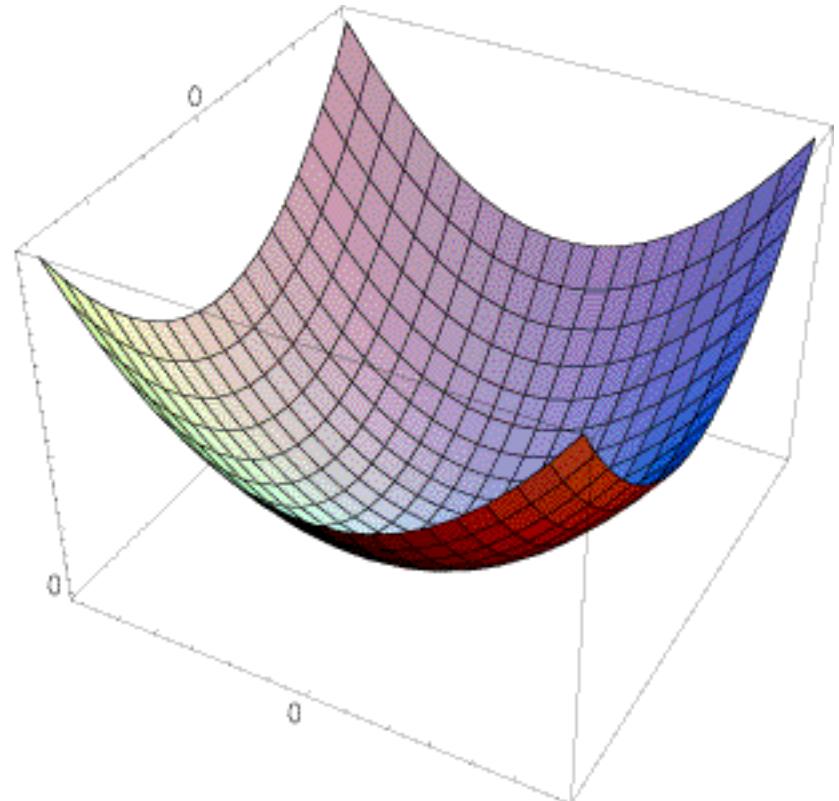
*What kind of image patch do these surfaces represent?*

# Visualization of a quadratic

The surface  $E(u,v)$  is locally approximated by a quadratic form

$$E(u,v) \approx [u \ v] H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$H = \sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

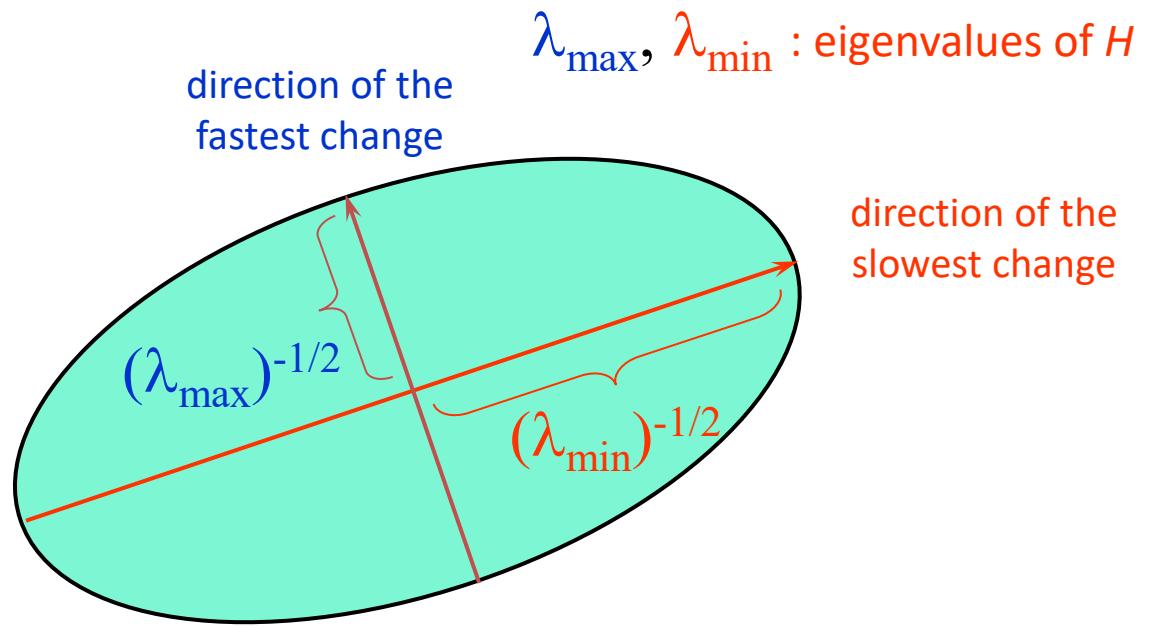


# General case

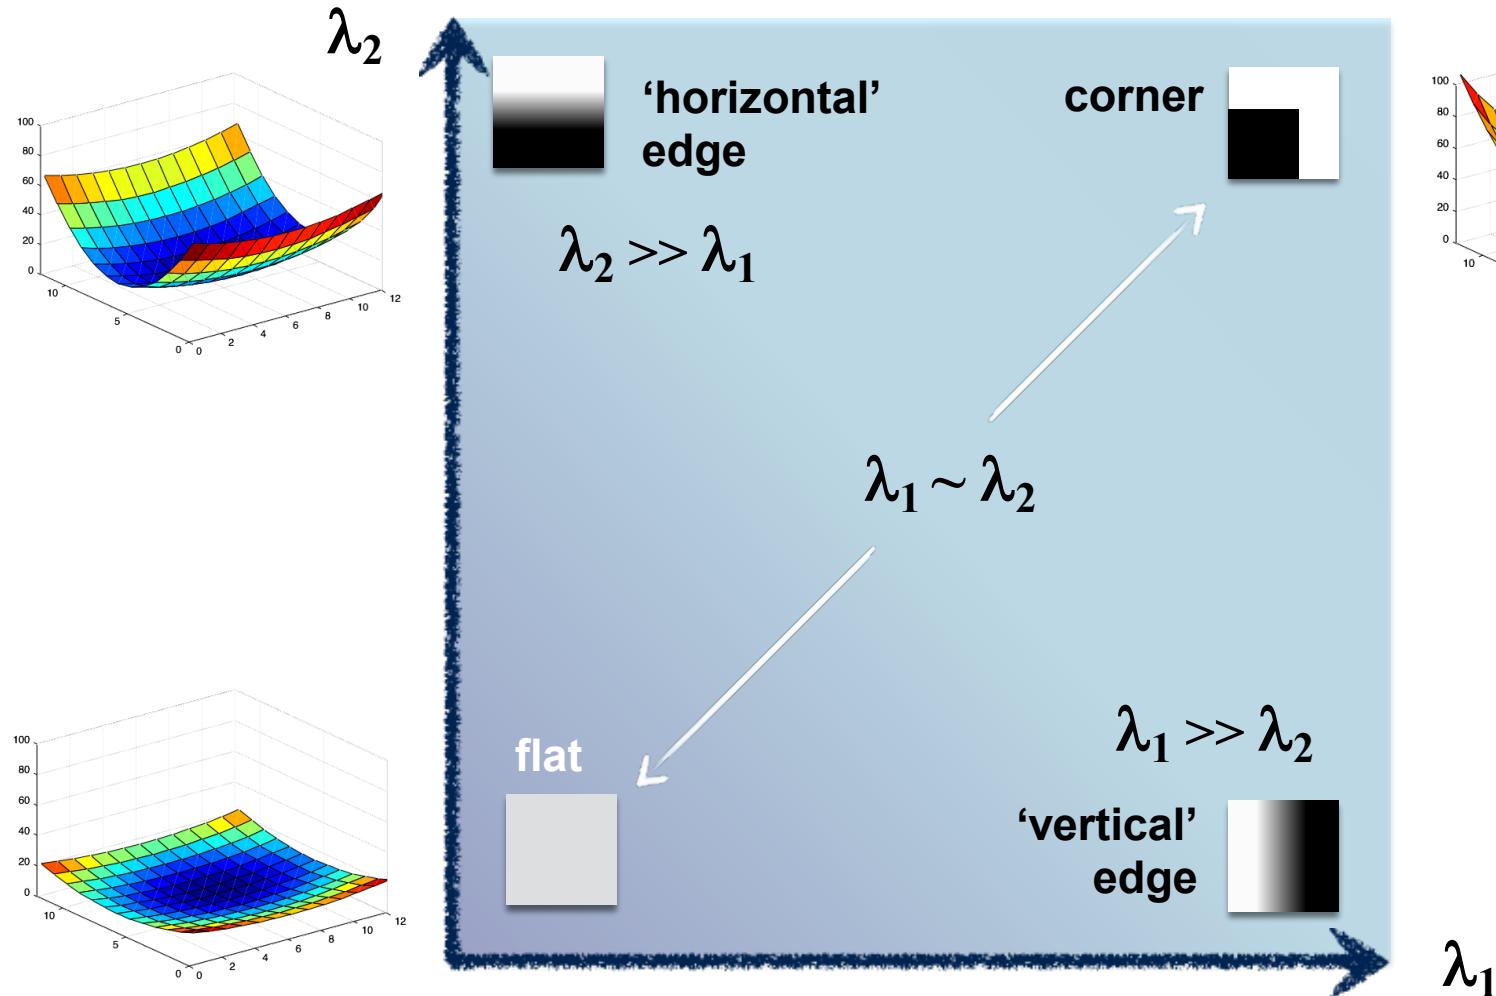
We can visualize  $H$  as an ellipse with axis lengths determined by the *eigenvalues* of  $H$  and orientation determined by the *eigenvectors* of  $H$

Ellipse equation:

$$[u \ v] H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



# interpreting eigenvalues



# Harris Detector [Harris88]

- Second moment matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives  
(optionally, blur first)

$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$
$$= \frac{\text{determinant}(H)}{\text{trace}(H)}$$

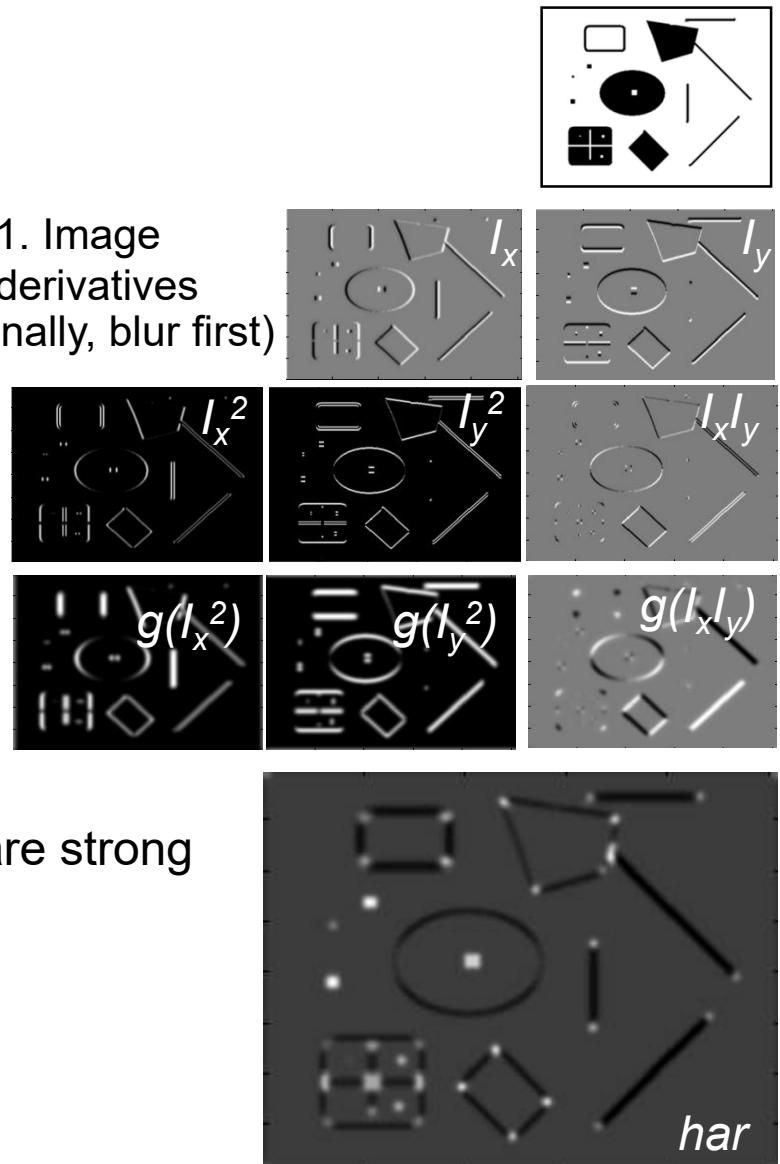
$$\det M = \lambda_1 \lambda_2$$
$$\text{trace } M = \lambda_1 + \lambda_2$$

2. Square of derivatives

3. Gaussian filter  $g(\sigma_f)$

4. Cornerness function – both eigenvalues are strong

5. Non-maxima suppression



# Harris features (in red)



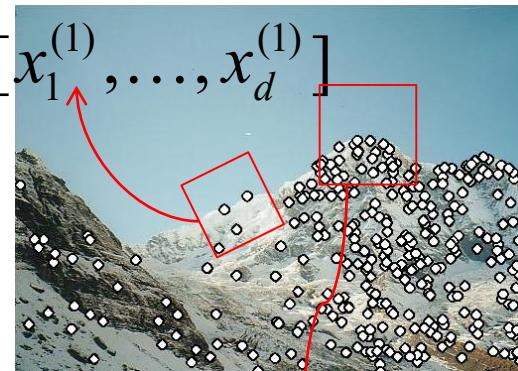
# Local features: main components

- 1) **Detection:** Identify the interest points

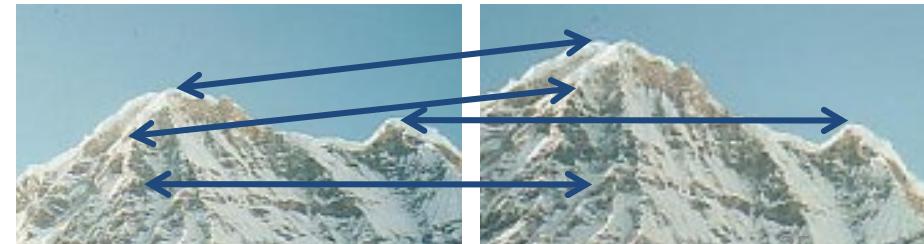
What is a good detector?

- 2) **Description:** Extract vector feature descriptor surrounding  $\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$  each interest point.

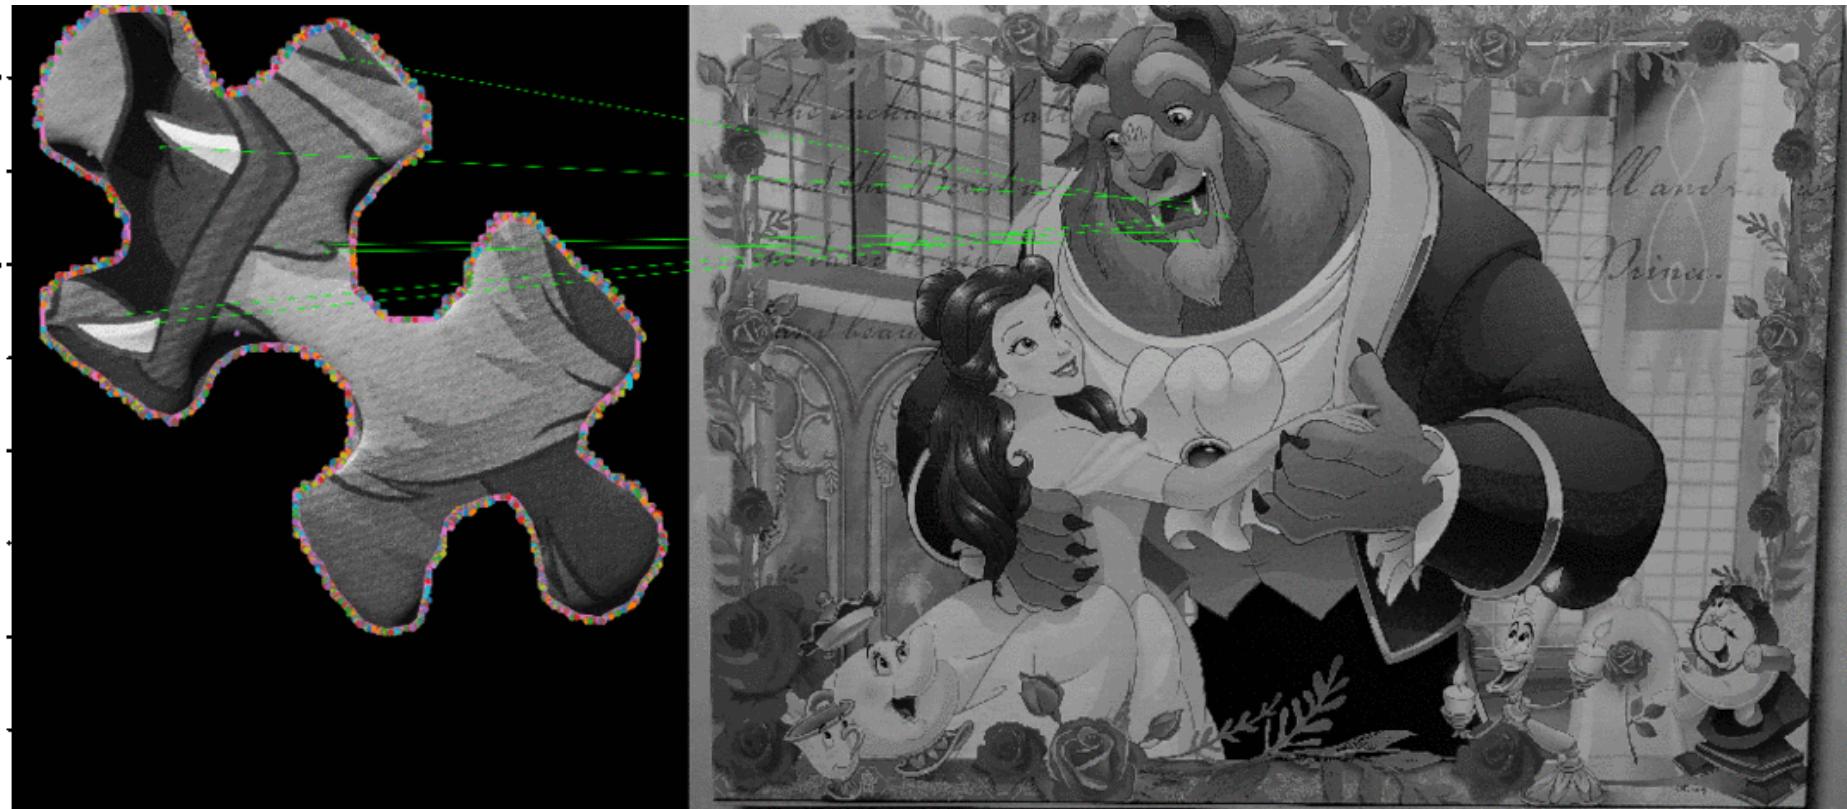
- 3) **Matching:** Determine correspondence between descriptors in two views



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



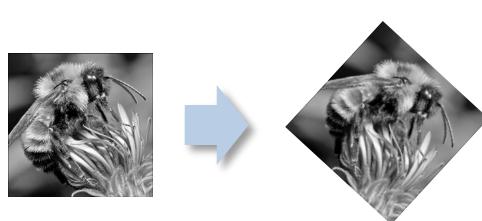
# Local Features



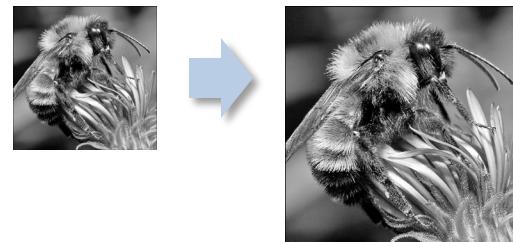
# Image transformations

- Geometric

**Rotation**



**Scale**



- Photometric

**Intensity change**



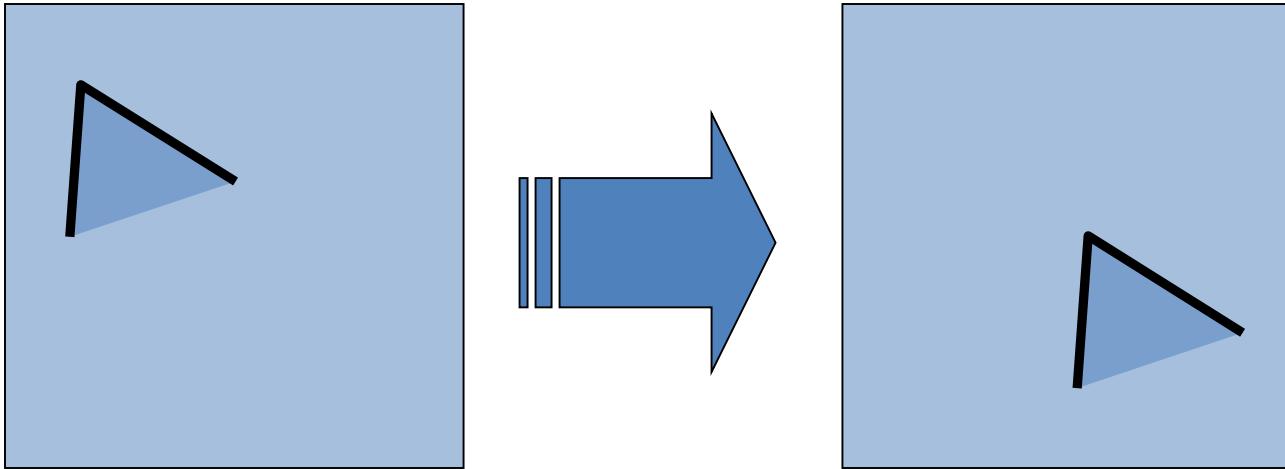
# Invariance and equivariance

- We want corner locations to be *invariant* to photometric transformations and *equivariant* to geometric transformations
  - **Invariance:** image is transformed and corner locations do not change
  - **Equivariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations
  - (Sometimes “invariant” and “equivariant” are both referred to as “invariant”)



# Harris detector: Invariance properties

## -- Image translation

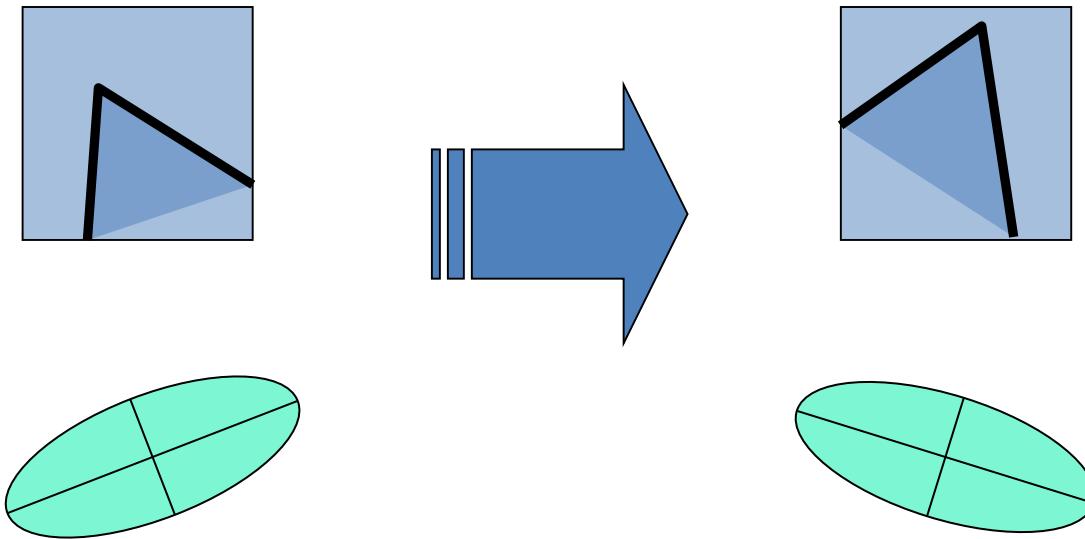


- Derivatives and window function are equivariant

Corner location is equivariant w.r.t. translation

# Harris detector: Invariance properties

## -- Image rotation



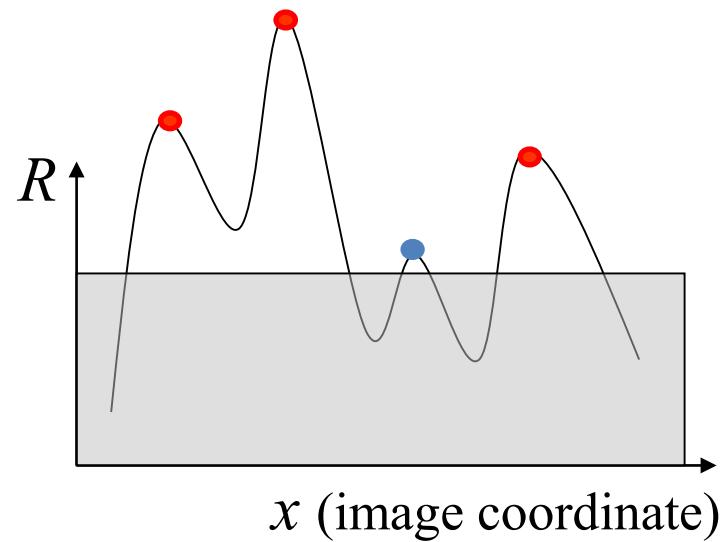
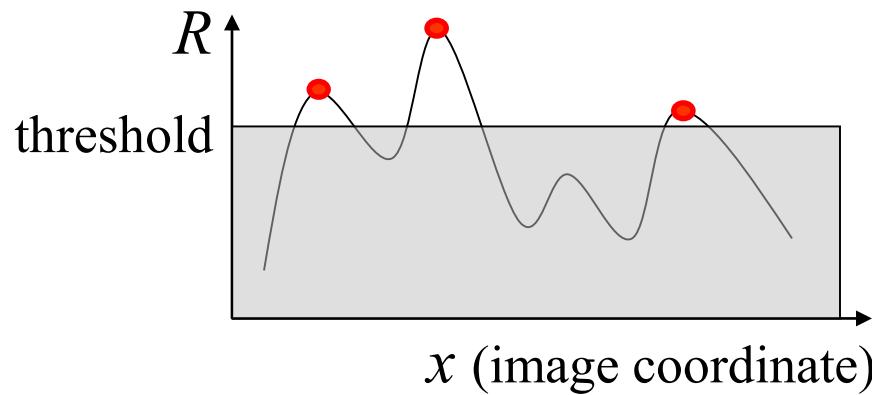
Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner location is equivariant w.r.t. image rotation

# Harris detector: Invariance properties – Affine intensity change

$$\begin{array}{c} \text{light blue square} \\ \xrightarrow{\quad} \\ \text{dark blue square} \end{array} \quad I \rightarrow a I + b$$

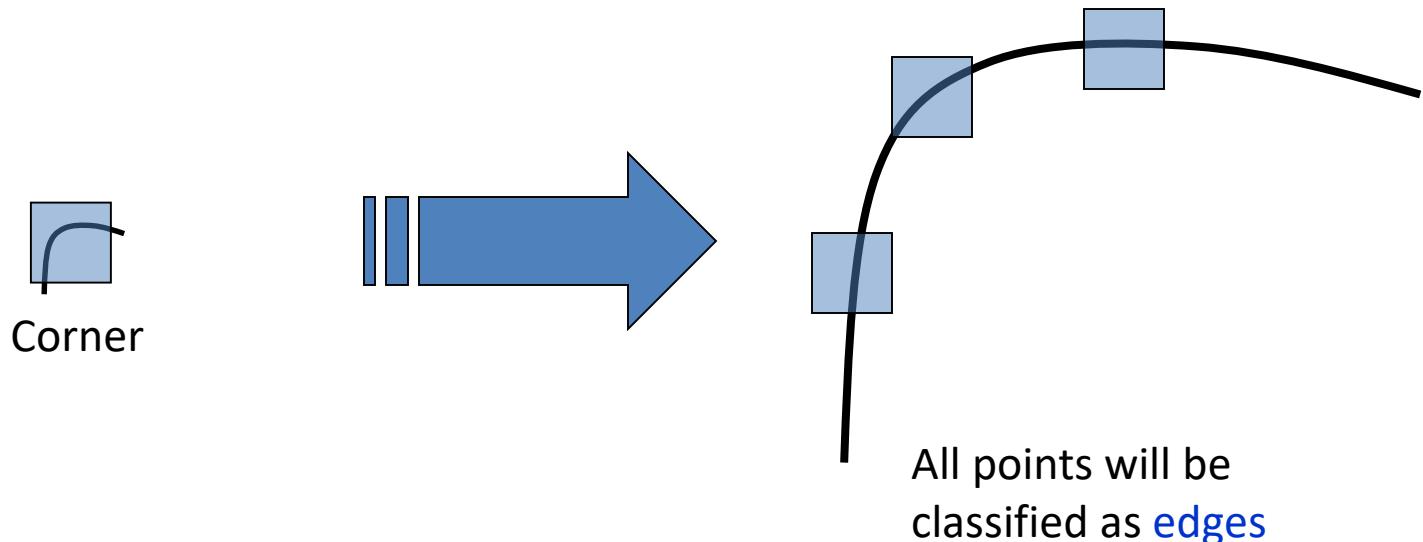
- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

# Harris Detector: Invariance Properties

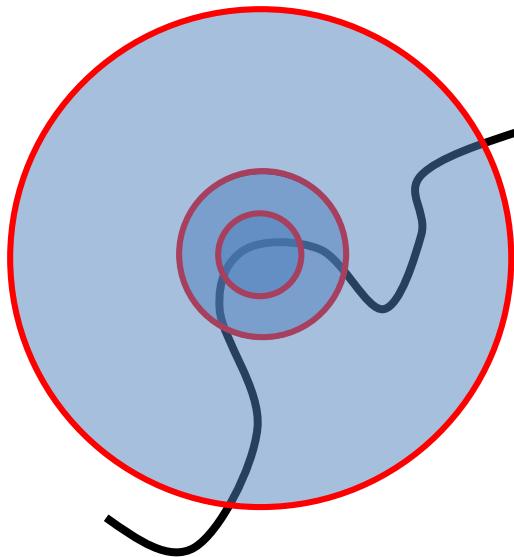
- Scaling



*Neither invariant nor equivariant to scaling*

# Scale invariant detection

Suppose you're looking for corners

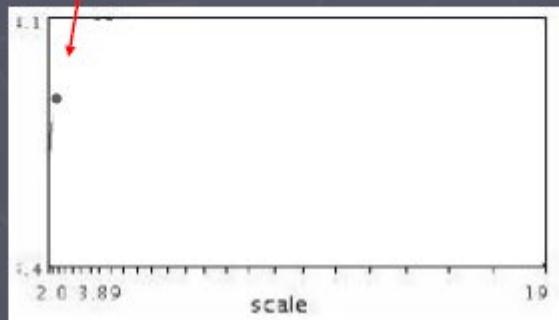


Key idea: find scale that gives local maximum of  $f$

- in both position and scale
- One definition of  $f$ : the Harris operator

# Automatic scale selection

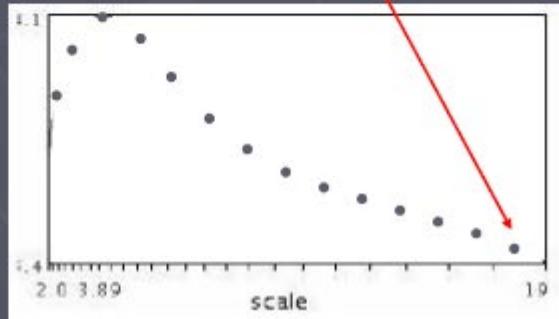
Lindeberg et al., 1996



$$f(I_{i_1\dots i_m}(x, \sigma))$$

Slide from Tinne Tuytelaars

# Automatic scale selection

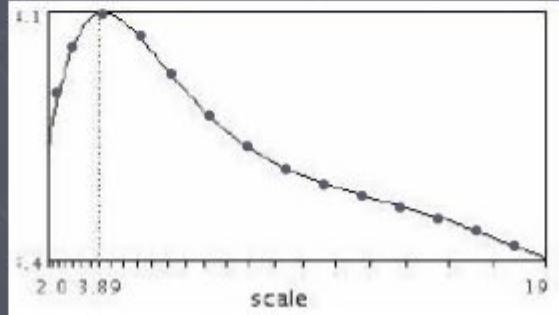


$f(I_{l_1\dots l_m}(x, \sigma))$

# Automatic scale selection

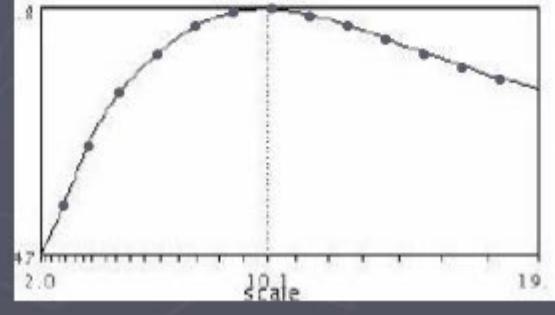


f



$f(I_{i_1 \dots i_m}(x, \sigma))$

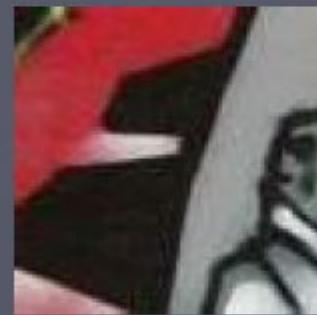
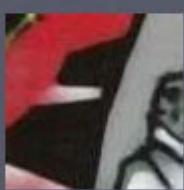
f



$f(I_{i_1 \dots i_m}(x', \sigma'))$

# Automatic scale selection

Normalize: rescale to fixed size



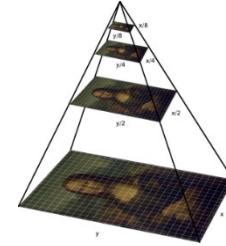
# Implementation

- Instead of computing  $f$  for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid

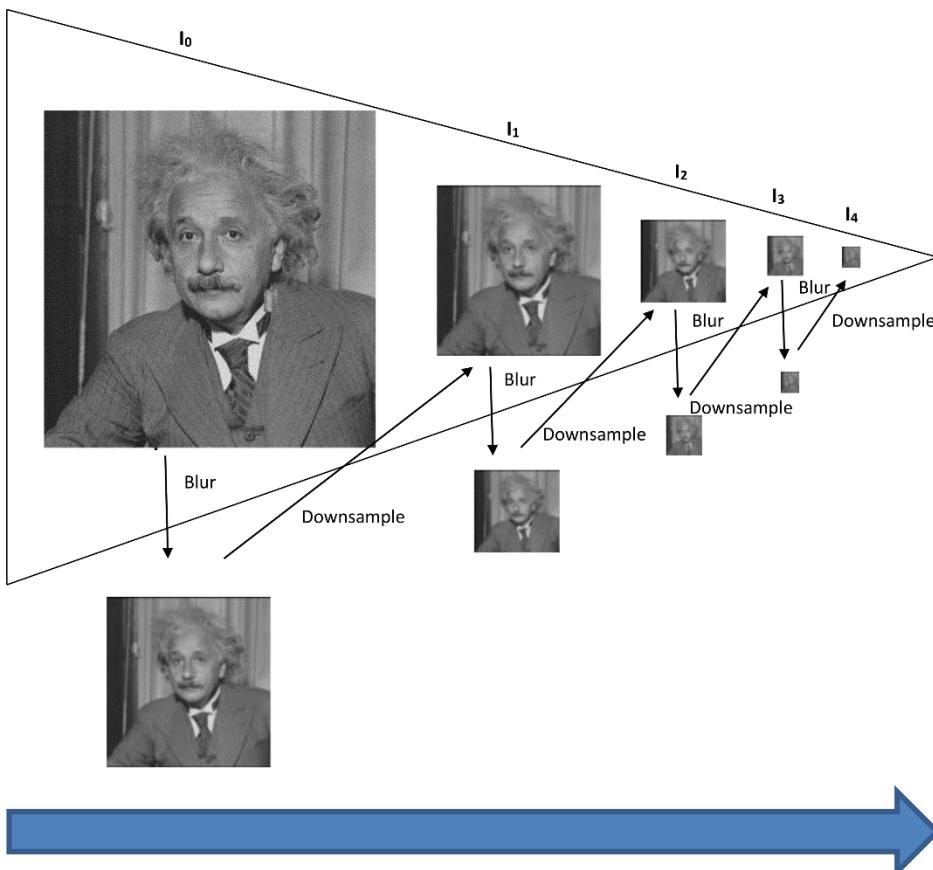


(sometimes need to create in-between levels, e.g. a  $\frac{3}{4}$ -size image)

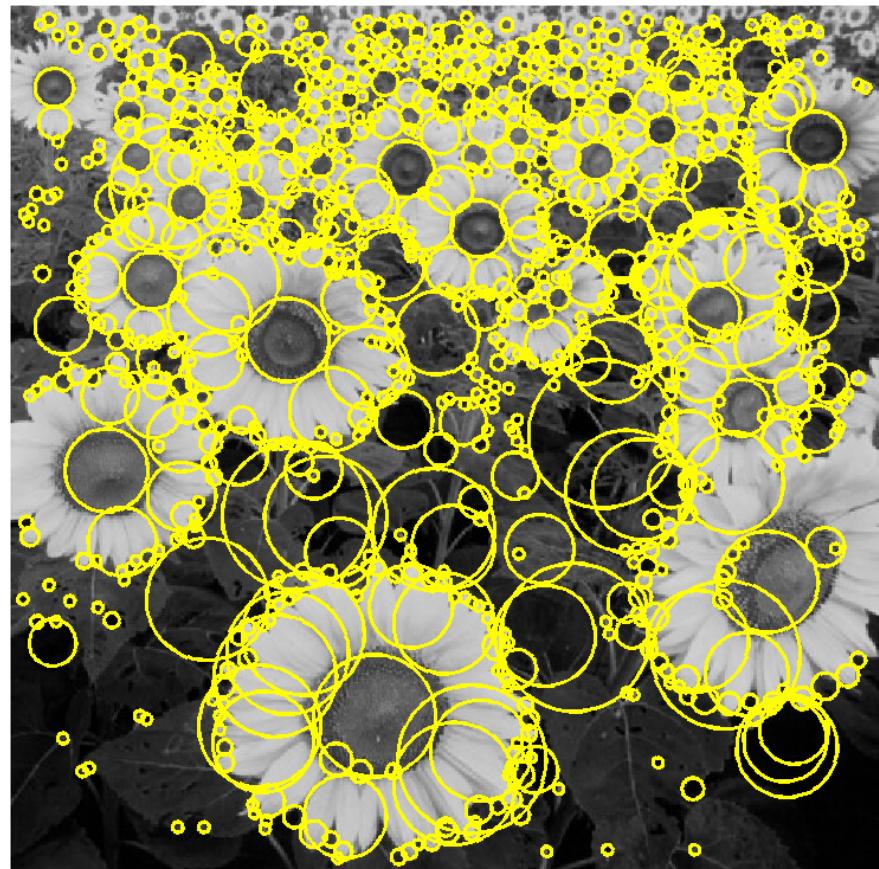
# Gaussian Pyramid?



- Blur -> Subsample -> Blur -> Subsample



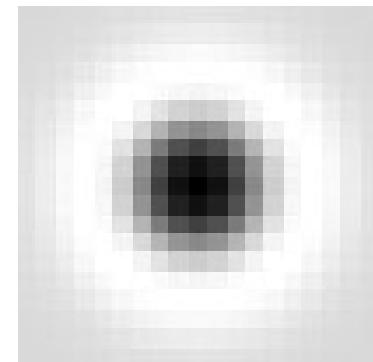
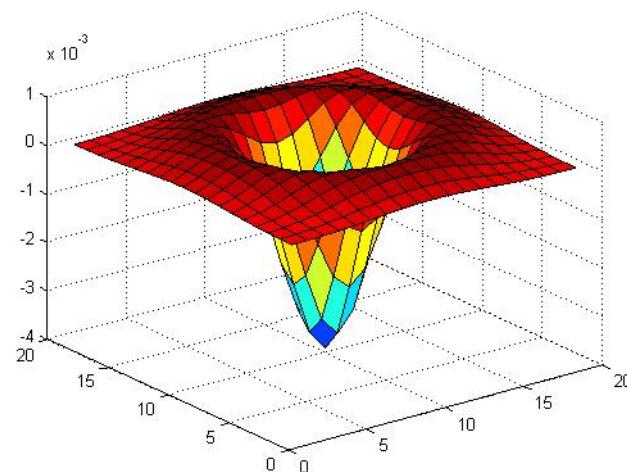
# Feature extraction: Corners and **blobs**





# Another common definition of $f$

- The *Laplacian of Gaussian (LoG)*

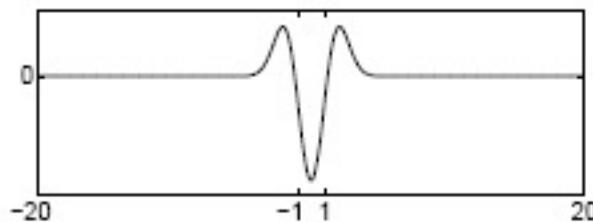


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

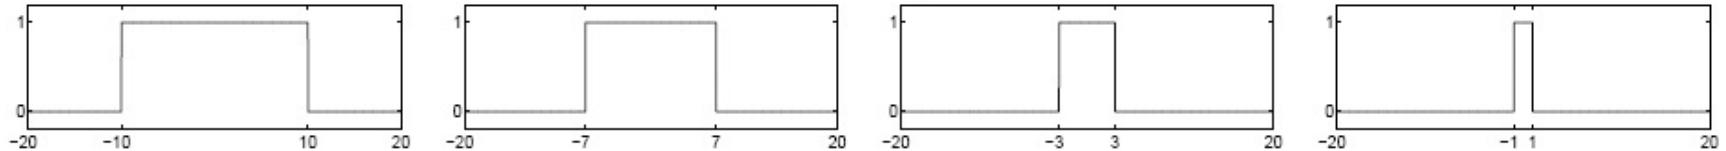
(very similar to a Difference of Gaussians (DoG) –  
i.e. a Gaussian minus a slightly smaller Gaussian)

Formally...

Laplacian filter



Original signal



Convolved with Laplacian ( $\sigma = 1$ )



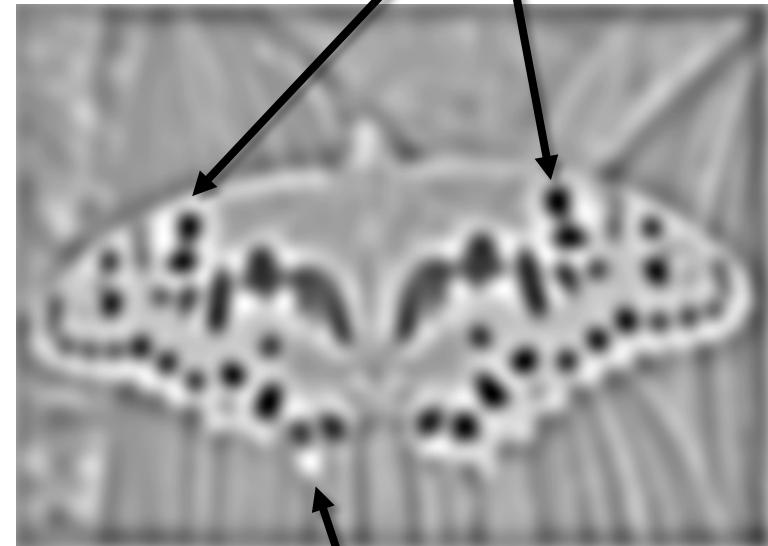
Highest response when the signal has the same **characteristic scale** as the filter

# Laplacian of Gaussian

- “Blob” detector



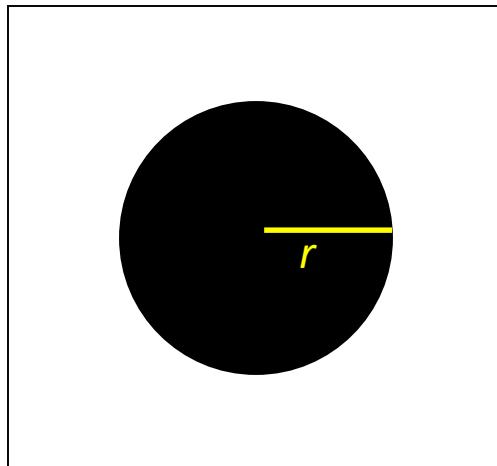
$$* \quad \bullet =$$



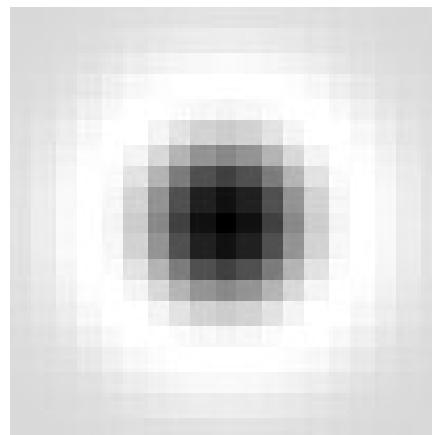
- Find maxima and minima of LoG operator in space and scale

# Scale selection

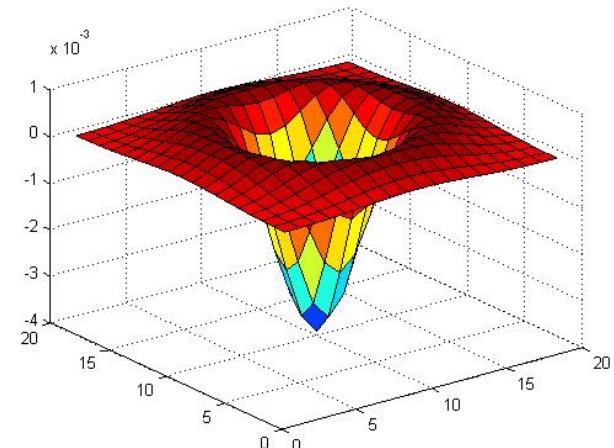
- At what scale does the Laplacian achieve a maximum response for a binary circle of radius  $r$ ?



image

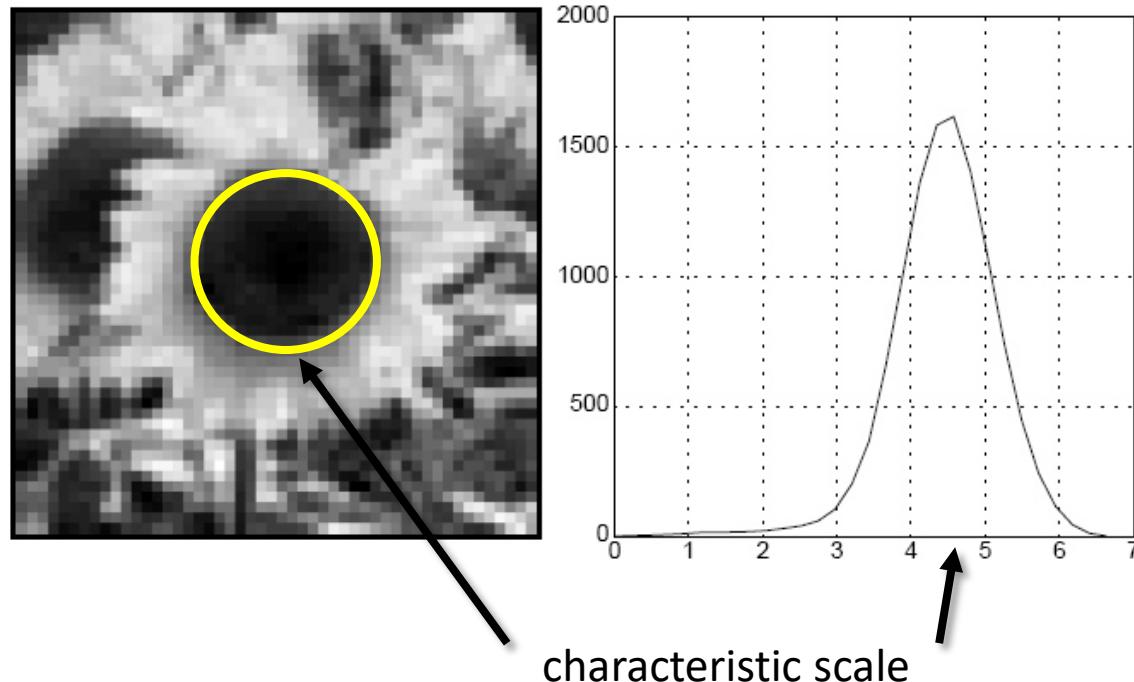


Laplacian



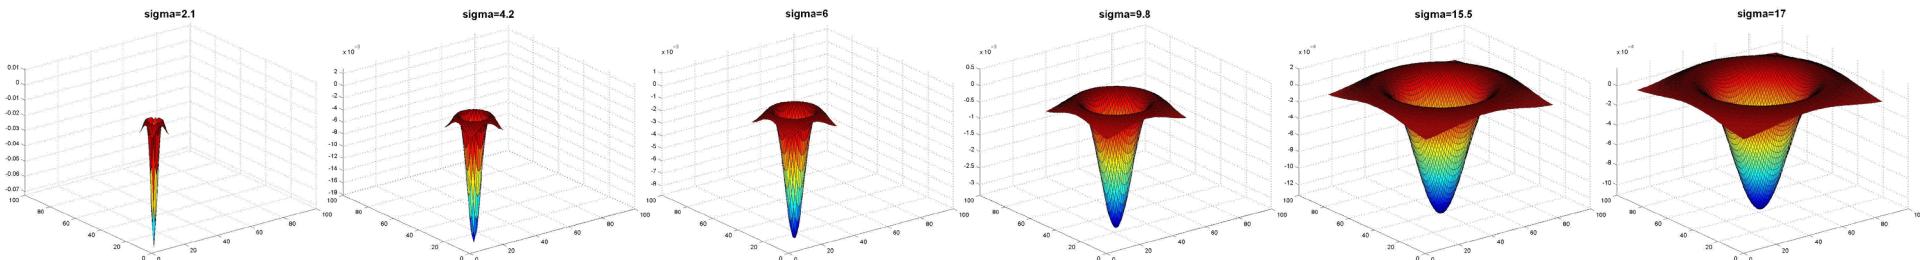
# Characteristic scale

- We define the characteristic scale as the scale that produces peak of Laplacian response



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)  
*International Journal of Computer Vision* **30** (2): pp 77--116.

# What happens if you apply different Laplacian filters?



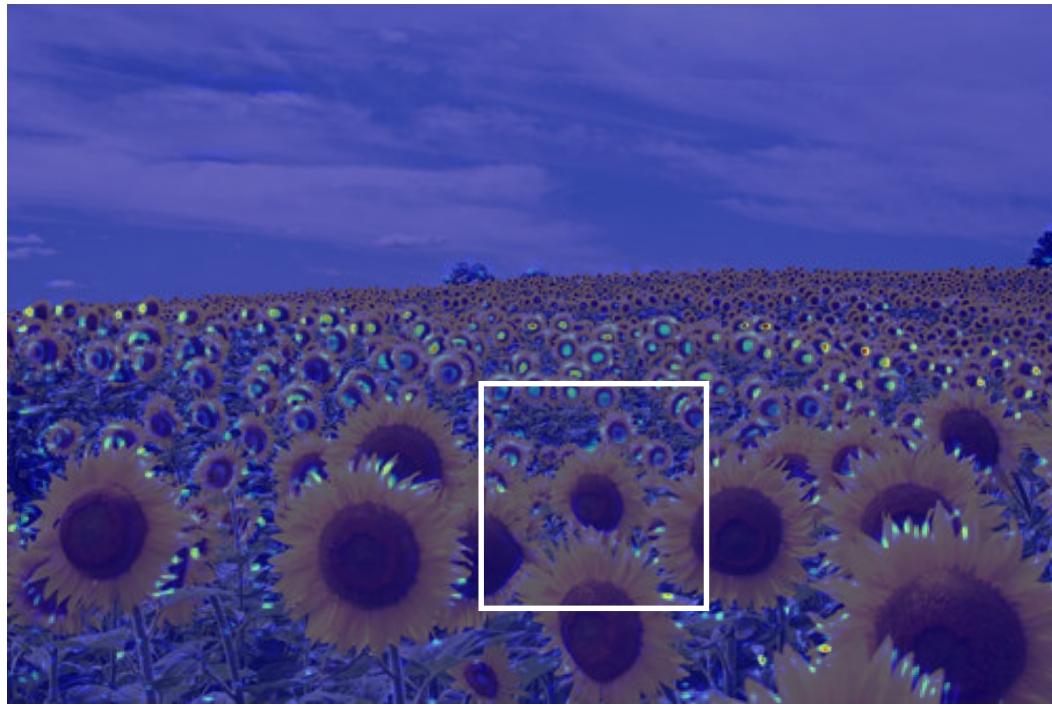
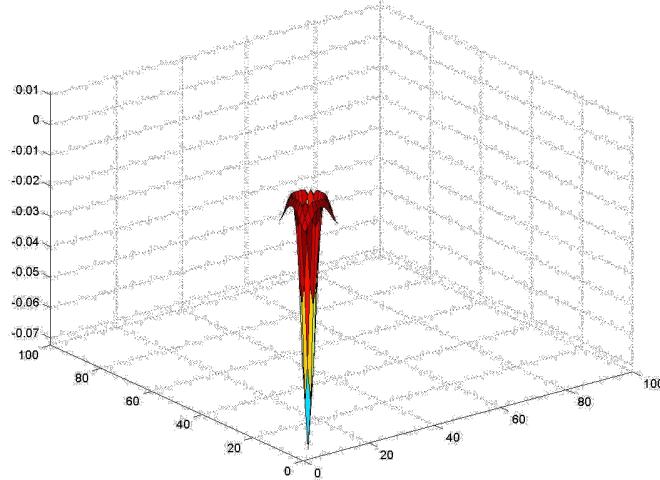
Full size

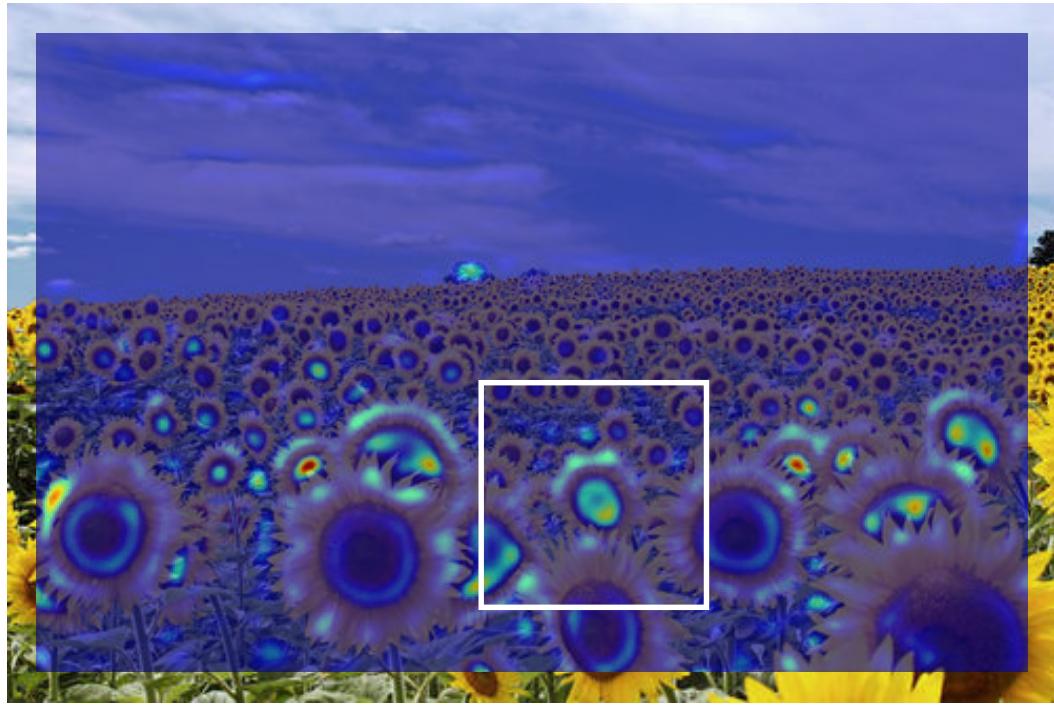
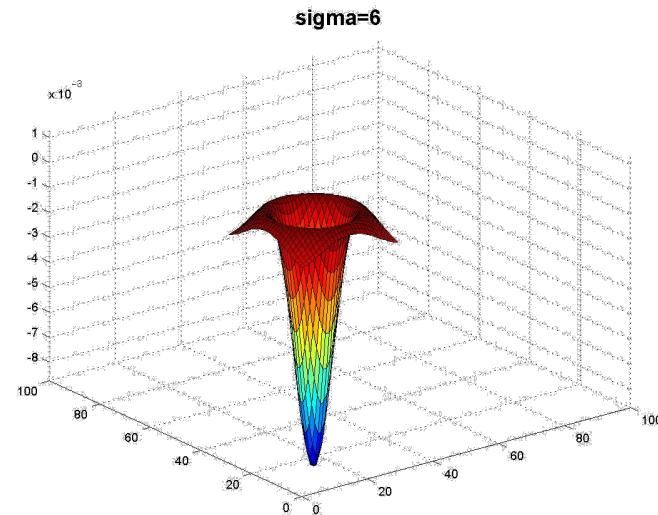


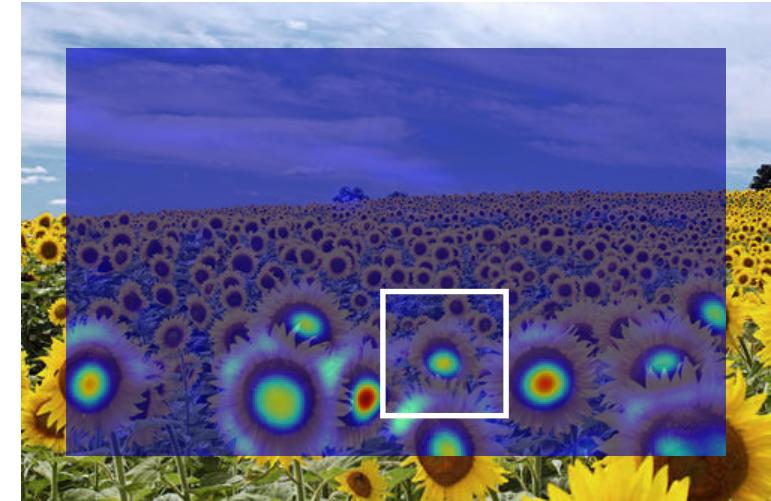
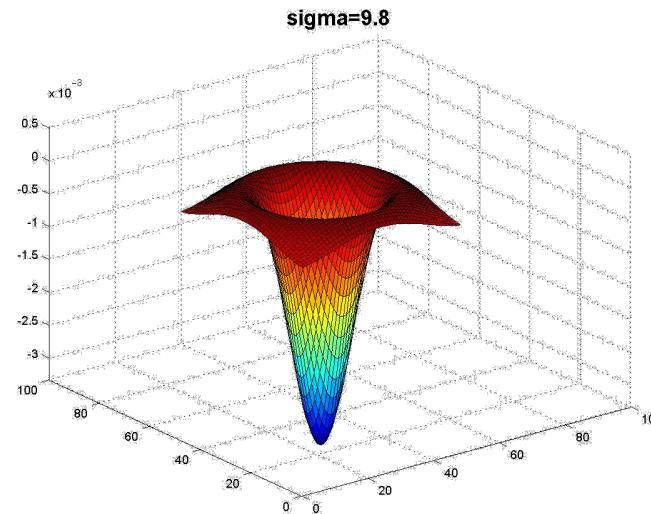
3/4 size



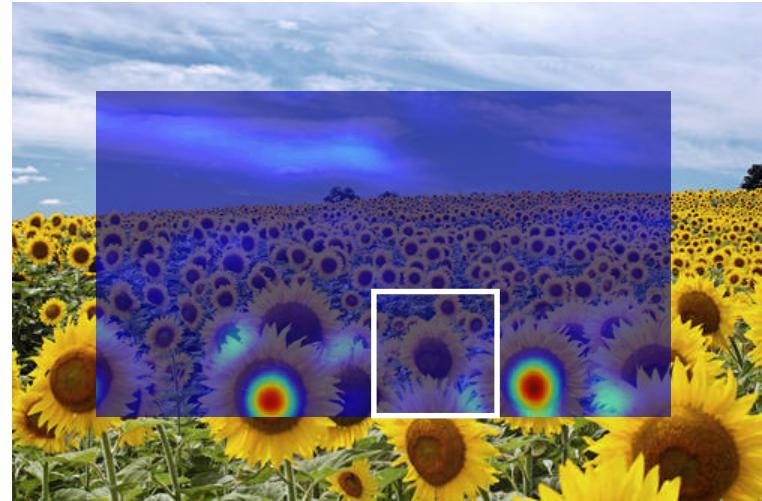
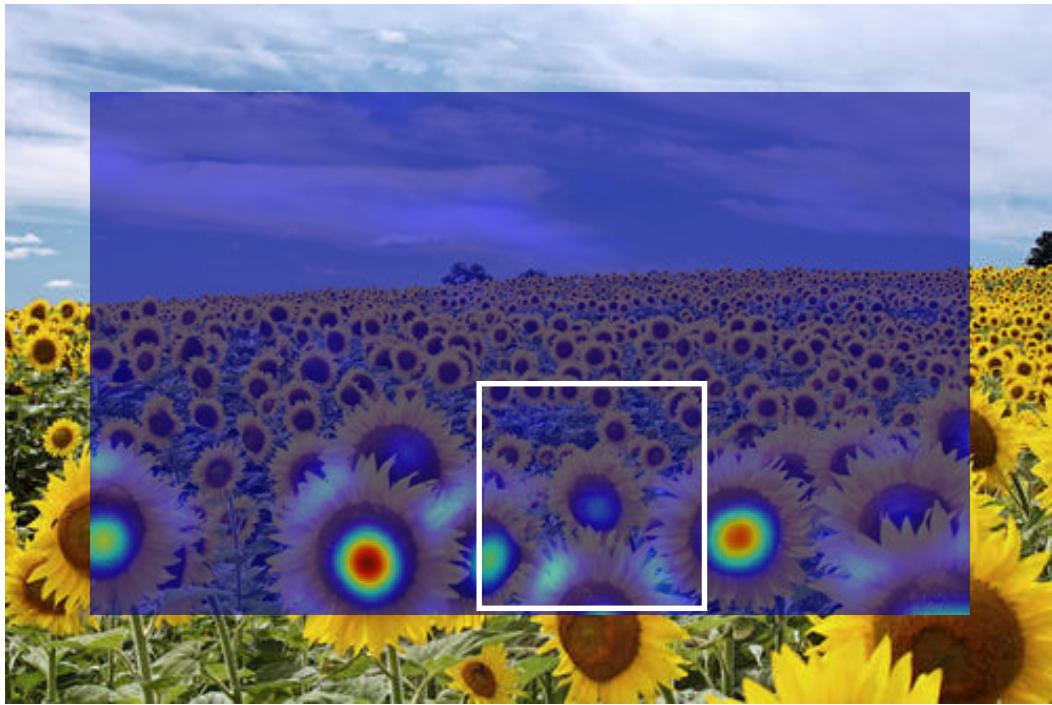
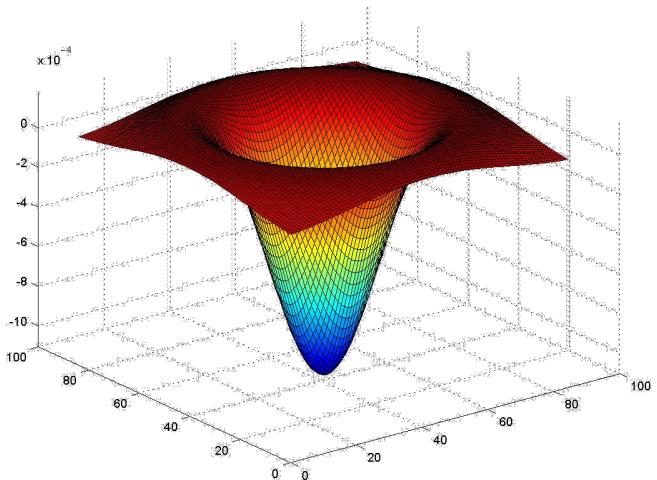
**sigma=2.1**







**sigma=17**



# Scale-space blob detector: Example

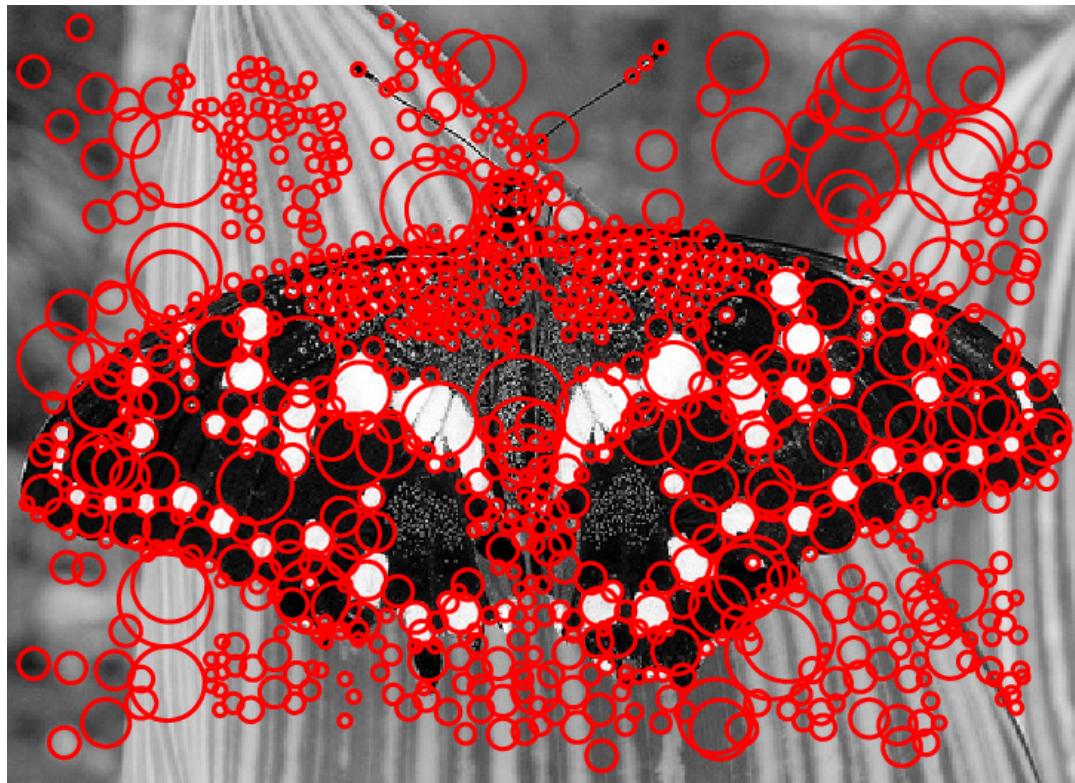


# Scale-space blob detector: Example



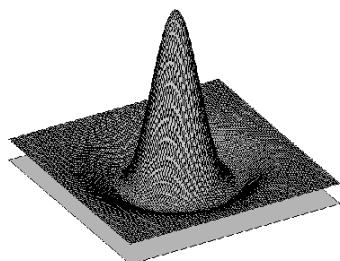
$\sigma = 11.9912$

# Scale-space blob detector: Example



# Find local maxima in 3D position-scale space

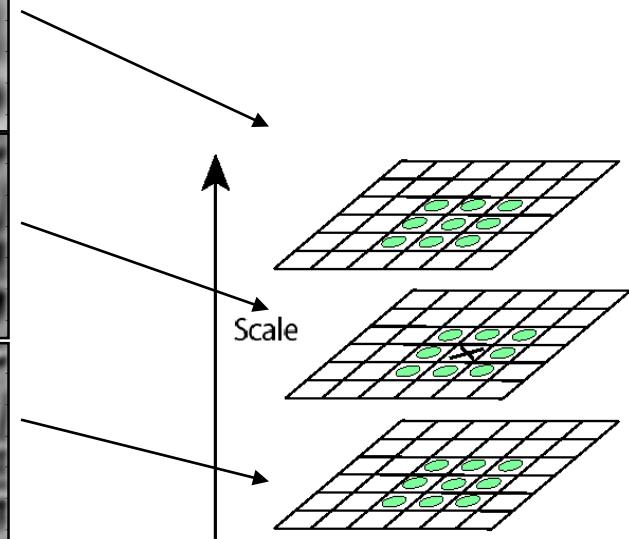
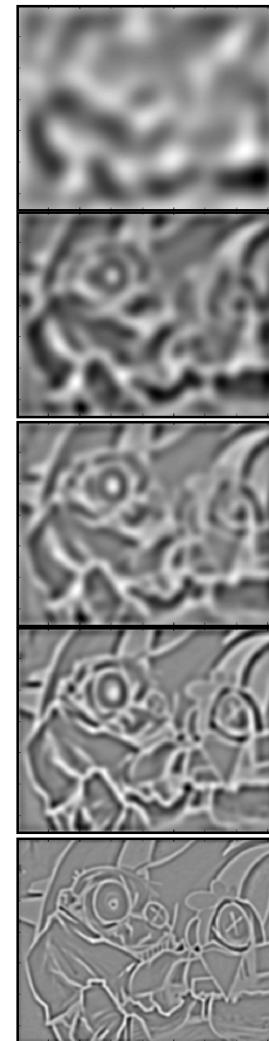
Advanced



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^3$$

Diagram illustrating the multi-scale convolution process:

- The input image is processed at different scales  $\sigma$ , represented by arrows pointing from the input to a series of five intermediate images.
- The intermediate images show increasing levels of blurring and edge detection, corresponding to scales  $\sigma^5, \sigma^4, \sigma^3, \sigma^2, \sigma$ .



⇒ List of  
 $(x, y, s)$

# Scale Invariant Detection

- Functions for determining scale

$$f = \text{Kernel} * \text{Image}$$

Kernels:

$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

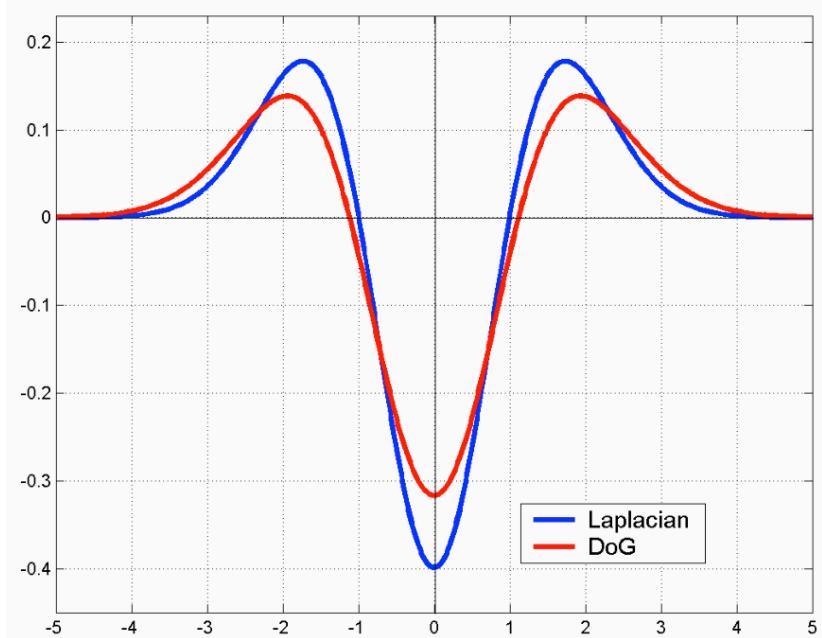
(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)

where Gaussian

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Note: The LoG and DoG operators are both rotation equivariant

# DoG with Scale-space

