

Object Detecion

<Vision System>




Department of Robot Engineering

Prof. Younggun Cho



Concept of Detection

- Given a category (ex. face, car, body), localizing objects in images.



Datasets



- Face detection
- One category: face
- Frontal faces
- Fairly rigid, unoccluded



1990's

Human Face Detection in Visual Scenes. H. Rowley, S. Baluja, T. Kanade. 1995.

Pedestrians



- One category: pedestrians
- Slight pose variations and small distortions
- Partial occlusions





Histograms of Oriented Gradients for Human Detection. N. Dalal and B. Triggs. CVPR 2005

PASCAL VOC




<http://host.robots.ox.ac.uk/pascal/VOC/>

visual object class


- 20 categories
- 10K images
- Large pose variations, heavy occlusions
- Generic scenes
- Cleaned up performance metric



Evaluation metric



Matching detections to ground truth



$$IoU(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

intersection over union

Matching detections to ground truth

- Match detection to most similar ground truth
 - highest IoU
- If $\text{IoU} > 50\%$, mark as correct
- If multiple detections map to same ground truth, mark only one as correct
- **Precision** = $\# \text{correct detections} / \text{total detections}$
- **Recall** = $\# \text{ground truth with matched detections} / \text{total ground truth}$

Tradeoff between precision and recall

- ML usually gives scores or probabilities, so threshold
- Too low threshold → too many detections
→ low precision, high recall
- Too high threshold → too few detections
→ high precision, low recall
- Right tradeoff depends on application
 - Detecting cancer cells in tissue: need high recall
 - Detecting edible mushrooms in forest: need high precision

Average precision



Average Precision (AP)



Average average precision


- AP marks detections with overlap $> 50\%$ as correct
- But may need better localization
- *Average* AP across multiple overlap thresholds
- Confusingly, still called average precision
- Introduced in COCO

Mean and category-wise AP

- Every category evaluated independently
- Typically report mean AP averaged over all categories
- Confusingly called “mean Average Precision”, or “**mAP**”

Why is detection hard(er)?

- Precise localization



Why is detection hard(er)?

- Much larger impact of pose




Why is detection hard(er)?

- Occlusion makes localization difficult



Why is detection hard(er)?

- Counting




Why is detection hard(er)?

- Small objects




Detection as classification

- Run through every possible box and classify
 - Well-localized object of class k or not?
- How many boxes?
 - Every pair of pixels = 1 box
 - $\binom{N}{2} = O(N^2)$
 - For 300 x 500 image, $N = 150K$
 - 2.25×10^{10} boxes!
- Related challenge: almost all boxes are negative!




Idea 1: scanning window

- Fix size
- Fix stride
- Crop boxes and classify
- Alternatively
 - Compute collection of feature maps
 - Convolve with filter




Multiple object sizes

- Objects can appear at any size
- *Discretize* set of sizes into a few different sizes
 - Sometimes called “anchors”
- Train separate classifier for each size




Dealing with large scale changes




Dealing with large scale changes

- Use an **image pyramid**
- Run same detector at multiple scales
- Take union of results



Idea 2: Object proposals

- Use segmentation to produce $\sim 5K$ “candidates”



Selective Search for Object Recognition

[J. R. R. Uijlings](#), [K. E. A. van de Sande](#), [T. Gevers](#), [A. W. M. Smeulders](#)

In International Journal of Computer Vision 2013.

Object proposals

- Basic idea: use grouping cues to identify segments that are likely to be objects
- Multiple versions
 - Do graph cuts with different seeds
 - Oversegment and then combinatorially group nearby objects

Example 1: Face detection



- Slides adapted Grauman & Liebe's tutorial
 - <http://www.vision.ee.ethz.ch/~bleibe/teaching/tutorial-aaai08/>
- Also see Paul Viola's talk (video)
 - <http://www.cs.washington.edu/education/courses/577/04sp/contents.html#DM>

Viola & Jones Face Detector: Rectangle filters




Why rectangles?


- Trick: *integral images* (aka *summed-area-tables*)

Integral images

What's the sum of pixels in the blue rectangle?




input image




integral image

Integral images


What's the sum of pixels in the rectangle?




A



B



C




D

Computing sum within a rectangle

- Let A,B,C,D be the values of the integral image at the corners of a rectangle
- Then the sum of original image values within the rectangle can be computed as:


$$\text{sum} = A - B - C + D$$

- **Only 3 additions** are required for any size of rectangle!




Filter as a classifier

How to convert the filter into a classifier?




Outputs of a rectangle
feature on faces and
non-faces.

Resulting weak classifier:


$$h_t(x) = \begin{cases} +1 & \text{if } f_t(x) > \theta_t \\ -1 & \text{otherwise} \end{cases}$$

Finding the best filters...




Considering all possible filter parameters: position, scale, and type:

180,000+ possible filters associated with each 24 x 24 window

Which of these filters(s) should we use to determine if a window has a face?


Boosting

**Weak
Classifier 1**




Boosting

**Weights
Increased**




Boosting




Boosting

**Weights
Increased**




Boosting



Boosting

**Final classifier is
a combination of weak
classifiers**



Boosting: training

- Initially, weight each training example equally
- In each boosting round:
 - find the weak classifier with lowest *weighted* training error
 - raise weights of training examples misclassified by current weak classifier
- Final classifier is linear combination of all weak classifiers
 - weight of each learner is directly proportional to its accuracy)
- Exact formulas for re-weighting and combining weak classifiers depend on the particular boosting scheme

- Given example images $(x_1, y_1), \dots, (x_n, y_n)$ where $y_i = 0, 1$ for negative and positive examples respectively.
- Initialize weights $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$ for $y_i = 0, 1$ respectively, where m and l are the number of negatives and positives respectively.
- For $t = 1, \dots, T$:

1. Normalize the weights,

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

so that w_t is a probability distribution.

- For each feature, j , train a classifier h_j which is restricted to using a single feature. The error is evaluated with respect to w_t , $\epsilon_j = \sum_i w_i |h_j(x_i) - y_i|$.
- Choose the classifier, h_t , with the lowest error ϵ_t .
- Update the weights:

$$w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$$

where $e_i = 0$ if example x_i is classified correctly, $e_i = 1$ otherwise, and $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$.


- The final strong classifier is:

$$h(x) = \begin{cases} 1 & \sum_{t=1}^T \alpha_t h_t(x) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

where $\alpha_t = \log \frac{1}{\beta_t}$

AdaBoost Algorithm

Start with
uniform weights
on training
examples



For T rounds


Evaluate
weighted error
for each feature,
pick best.

Re-weight the examples:
Incorrectly classified -> more weight
Correctly classified -> less weight

Final classifier is combination of the
weak ones, weighted according to
error they had.

Freund & Schapire 1995

Viola-Jones Face Detector: Results

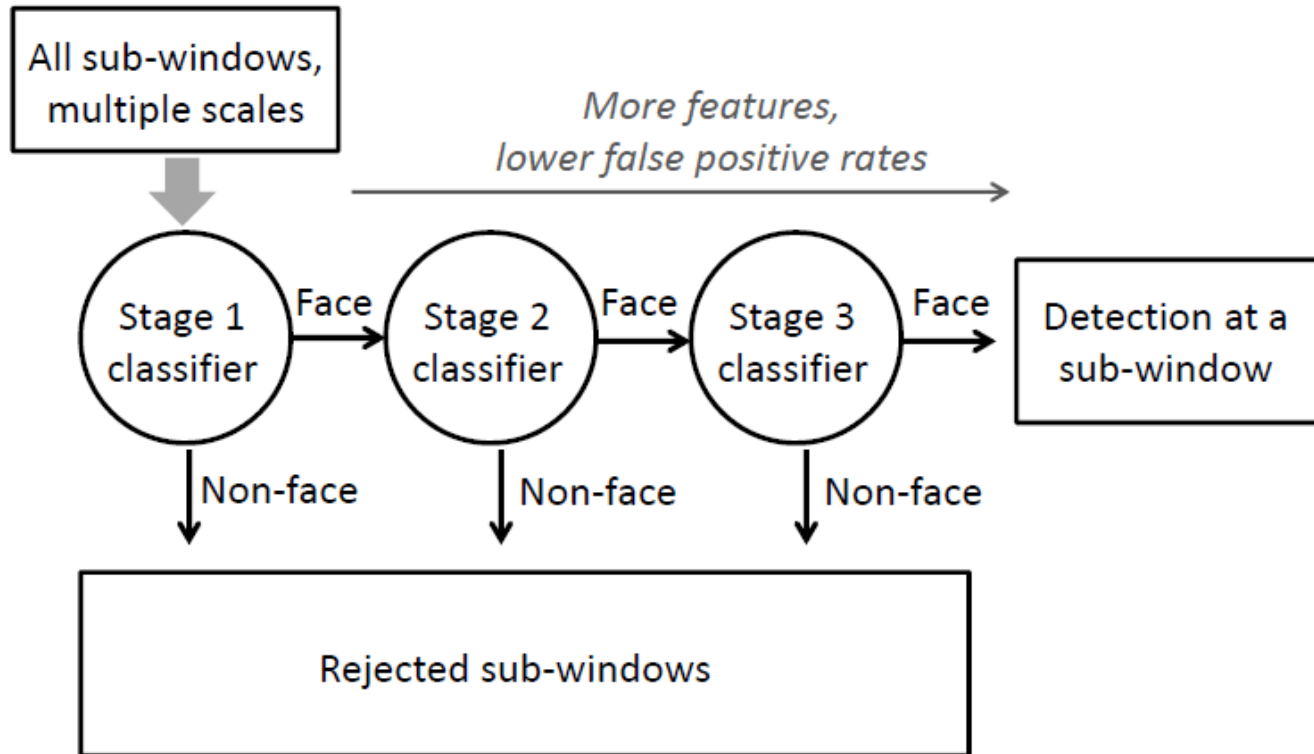


First two features
selected

Robust Real-Time Face Detection ,IJCV, 2004

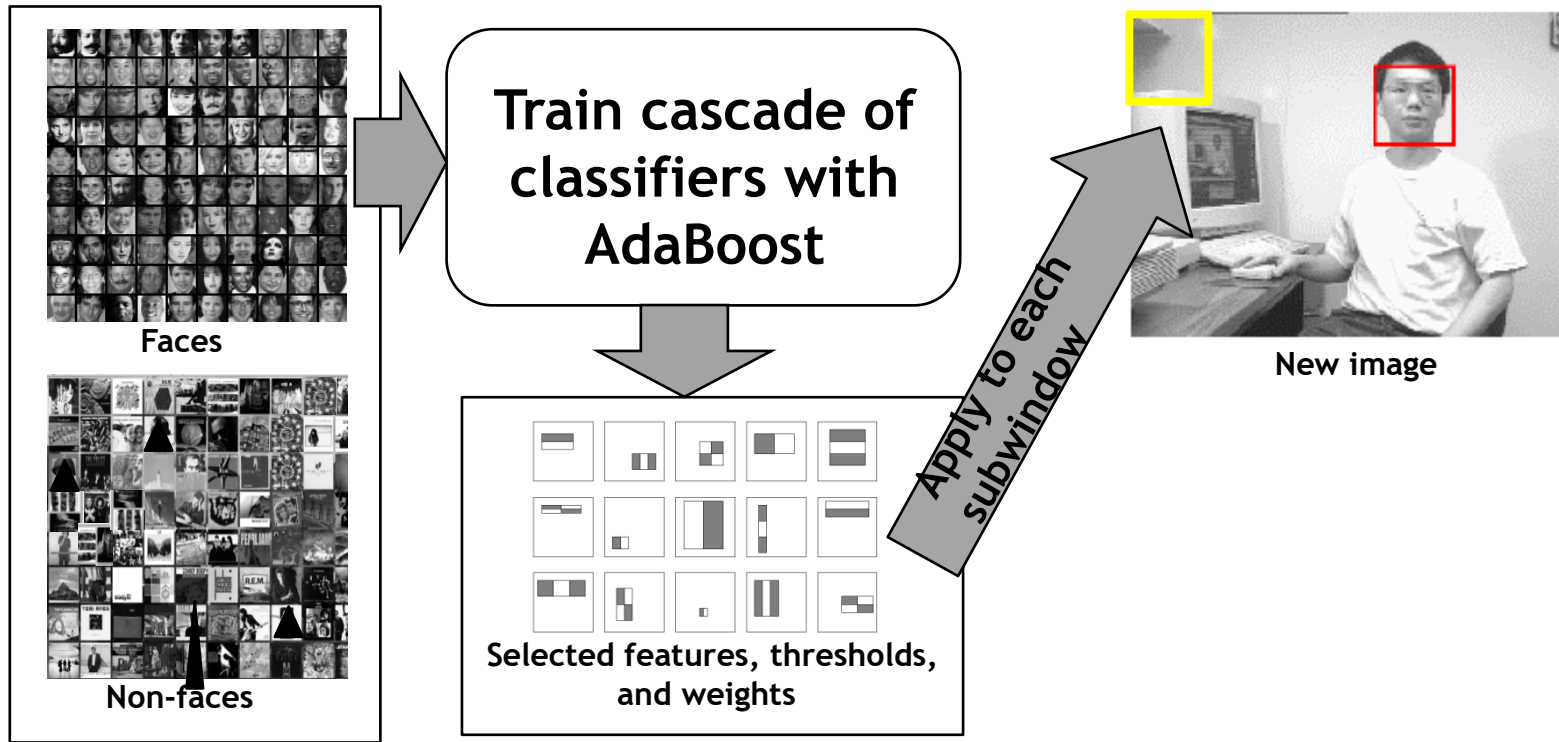
- Even if the filters are fast to compute, each new image has a lot of possible windows to search.
- How to make the detection more efficient?

Cascading classifiers for detection



- Form a *cascade* with low false negative rates early on
- Apply less accurate but faster classifiers first to immediately discard windows that clearly appear to be negative





Viola-Jones detector: summary



- Train with 5K positives, 350M negatives
- Real-time detector using 38 layer cascade
- 6061 features in all layers


• [Implementation available in OpenCV: <http://www.intel.com/technology/computing/opencv/>]

Viola-Jones Face Detector: Results



Example 2: Pedestrian detection

- Detecting upright, walking humans also possible using sliding window's appearance/texture; e.g.,






SVM with HoG [Dalal & Triggs,
CVPR 2005]

Dalal-Triggs pedestrian detector



1. Extract fixed-sized (64x128 pixel) window at each position and scale
2. Compute HOG (histogram of gradient) features within each window
3. Score the window with a linear SVM classifier
4. Perform non-maxima suppression to remove overlapping detections with lower scores





- **Tested with**


- RGB
- LAB
- Grayscale

} Slightly better performance vs. grayscale

- **Gamma Normalization and Compression**

- Square root
- Log

} Very slightly better performance vs. no adjustment



Outperforms



-1	0	1
----	---	---


centered

-1	1
----	---

uncentered

1	-8	0	8	-1
---	----	---	---	----

cubic-corrected




0	1
-1	0

diagonal


-1	0	1
-2	0	2
-1	0	1

Sobel




- **Histogram of gradient orientations**


Orientation: 9 bins (for unsigned angles)



Histograms in $k \times k$ pixel cells




- Votes weighted by magnitude
- Bilinear interpolation between cells






R-HOG



Normalize with respect to surrounding cells

$$L2 - norm : v \longrightarrow v / \sqrt{\|v\|_2^2 + \epsilon^2}$$



Original Formulation


orientations



$$\# \text{ features} = 15 \times 7 \times 9 \times 4 = 3780$$

cells


normalizations by
neighboring cells



X=





pos w neg w





$$0.16 = w^T x - b$$

$$\text{sign}(0.16) = 1$$

\Rightarrow pedestrian

Detection examples



Something to think about...

- Sliding window detectors work
 - *very well* for faces
 - *fairly well* for cars and pedestrians
 - *badly* for cats and dogs
- Why are some classes easier than others?

Strengths and Weaknesses of Statistical Template Approach

Strengths


- Works very well for non-deformable objects with canonical orientations: faces, cars, pedestrians
- Fast detection

Weaknesses

- Not so well for highly deformable objects or “stuff”
- Not robust to occlusion
- Requires lots of training data

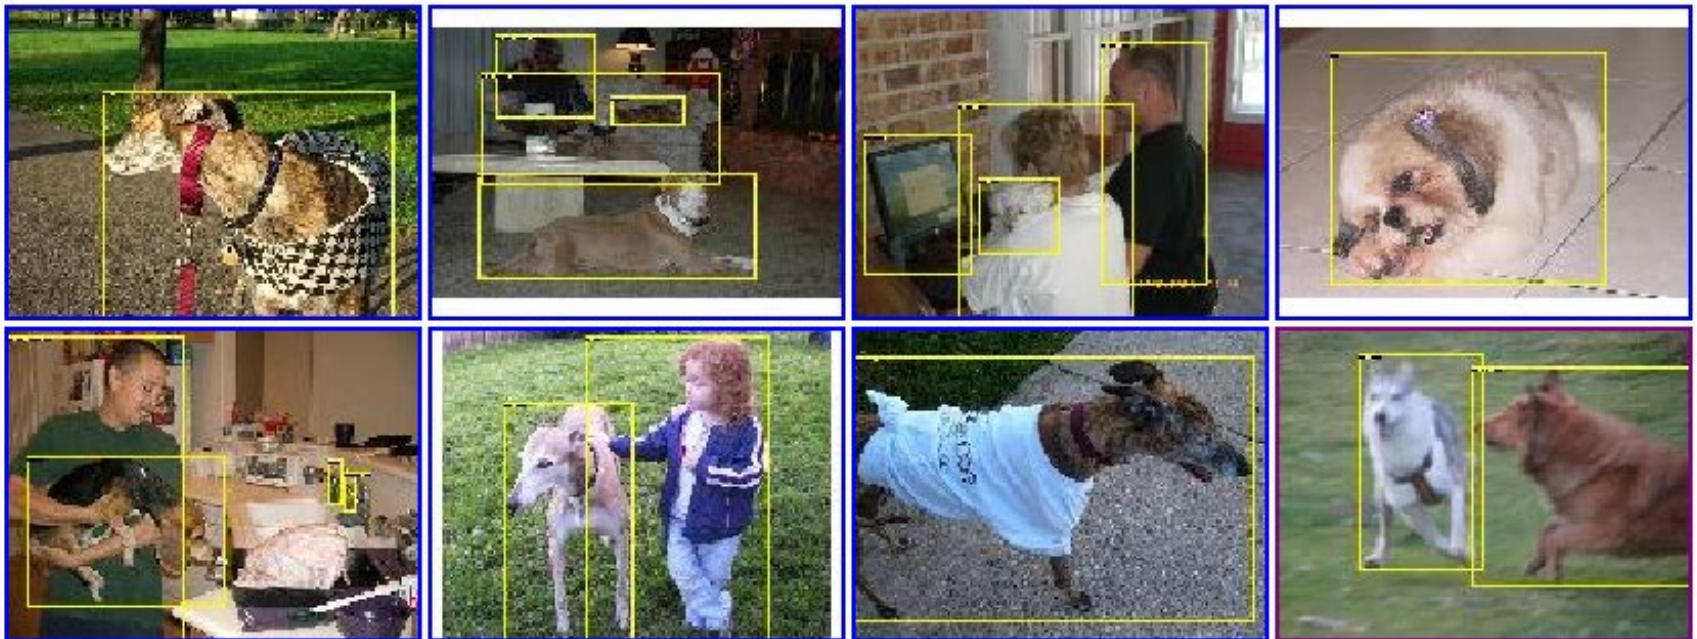
Limitations (continued)

- Not all objects are “box” shaped



Limitations (continued)

- Non-rigid, deformable objects not captured well with representations assuming a fixed 2d structure; or must assume fixed viewpoint
- Objects with less-regular textures not captured well with holistic appearance-based descriptions



Limitations (continued)

- If considering windows in isolation, context is lost



Sliding window



Detector's view

Limitations (continued)

- In practice, often entails large, cropped training set (expensive)
- Requiring good match to a global appearance description can lead to sensitivity to partial occlusions



Tricks of the trade




- **Details in feature computation really matter**
 - E.g., normalization in Dalal-Triggs improves detection rate by 27% at fixed false positive rate
- **Template size**
 - Typical choice is size of smallest detectable object
- **“Jittering” to create synthetic positive examples**
 - Create slightly rotated, translated, scaled, mirrored versions as extra positive examples
- **Bootstrapping to get hard negative examples**
 1. Randomly sample negative examples
 2. Train detector
 3. Sample negative examples that score > -1
 4. Repeat until all high-scoring negative examples fit in memory

Influential Works in Detection


- **Sung-Poggio (1994, 1998) : ~2000 citations**
 - Basic idea of statistical template detection (I think), bootstrapping to get “face-like” negative examples, multiple whole-face prototypes (in 1994)
- **Rowley-Baluja-Kanade (1996-1998) : ~3600**
 - “Parts” at fixed position, non-maxima suppression, simple cascade, rotation, pretty good accuracy, fast
- **Schneiderman-Kanade (1998-2000,2004) : ~1700**
 - Careful feature engineering, excellent results, cascade
- **Viola-Jones (2001, 2004) : ~11,000**
 - Haar-like features, Adaboost as feature selection, hyper-cascade, very fast, easy to implement
- **Dalal-Triggs (2005) : ~6500**
 - Careful feature engineering, excellent results, HOG feature, online code
- **Felzenszwalb-Huttenlocher (2000): ~2100**
 - Efficient way to solve part-based detectors
- **Felzenszwalb-McAllester-Ramanan (2008): ~1300**
 - Excellent template/parts-based blend

Things to remember

- Sliding window for search
- Features based on differences of intensity (gradient, wavelet, etc.)
 - Excellent results require careful feature design
- Boosting for feature selection
- Integral images, cascade for speed
- Bootstrapping to deal with many, many negative examples



Lab. 1 Face Detection



Code

<https://github.com/opencv/opencv/tree/master/data/haarcascades>

Download here



```
1 import cv2
2
3 # Load the cascade
4 face_cascade = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
5 # Read the input image
6 img = cv2.imread('Lena.jpg')
7 # Convert into grayscale
8 gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
9 # Detect faces
10 faces = face_cascade.detectMultiScale(gray, 1.1, 4)
11 # Draw rectangle around the faces
12 for (x, y, w, h) in faces:
13     cv2.rectangle(img, (x, y), (x+w, y+h), (255, 0, 0), 2)
14 # Display the output
15 cv2.imshow('img', img)
16 cv2.waitKey()
```